

PeopleFlow: A Reproducible Geometry-Aware Multi-Agent Framework for Evacuation Safety Evaluation in Smart Building Environments

Likith Sai Kunchi^{1,*}, Manivannan Thulasirangan Lakshmidēvi², K S K Karthik Reddy³

^{1,*}*School of Computer Science and Engineering (SCOPE), VIT (Vellore Institute of Technology), Vellore, Tamil Nadu, India*

²*School of Civil Engineering, VIT (Vellore Institute of Technology), Vellore, Tamil Nadu, India*

³*School of Civil Engineering, VIT (Vellore Institute of Technology), Vellore, Tamil Nadu, India*

Correspondence*:

Likith Sai Kunchi

likithsai.kunchi@gmail.com

2 ABSTRACT

3 Evaluating evacuation performance in complex buildings requires analysis methods that capture
4 interactions between spatial configuration, pedestrian behavior, and operational disruptions.
5 This paper introduces PeopleFlow, a layout-informed computational framework for structured
6 evacuation-safety comparison across diverse architectural layouts. The workflow converts floor-
7 plan representations into navigable simulation environments and applies continuous-space
8 multi-agent modeling based on force-driven pedestrian dynamics. PeopleFlow reports inter-
9 pretable indicators including total clearance duration, localized congestion intensity, bottleneck
10 persistence, and exit-utilization balance. To enable consistent experimentation, the framework
11 uses parameterized scenario configurations and reproducibility-oriented execution records. The
12 evaluation covers a structured matrix of nine layout types across five operational conditions (450
13 core runs), plus 210 controlled supplementary/sensitivity runs for real-layout generalization and
14 sensitivity to exit width, desired speed, and occupancy scaling. Results show strong depende-
15 nce of evacuation outcomes on geometric topology and routing strategy, with clearance-time
16 differences exceeding an eightfold range and substantial congestion variation across layouts.
17 The primary contribution is a reproducible engineering framework for relative safety prioritization
18 across heterogeneous architectural configurations, rather than deterministic prediction of indivi-
19 dual trajectories. The framework supports what-if analysis during early design and provides a
20 computational basis for integration with digital twin workflows for high-occupancy safety planning.
21 From a built-environment perspective, the framework enables early-stage design evaluation of
22 evacuation efficiency and supports decision-making in smart building planning, infrastructure
23 safety assessment, and digital-twin-based emergency management systems.

24 **Keywords:** Crowd simulation, evacuation analytics, evacuation modelling, digital twin, multi-agent modelling, multi-agent simulation,
25 smart infrastructure, building evacuation safety, reproducible pipeline.

1 INTRODUCTION

Building evacuation performance evaluation is a high-impact engineering task where static compliance checks often miss dynamic crowd effects under disruption. Historical emergencies consistently show that bottleneck formation, uneven exit utilization, and localized panic response can dominate outcomes even in code-compliant facilities. Conventional hand calculations remain useful for baseline checks, but they provide limited support for comparing design alternatives under multi-hazard and operational uncertainty Hurley et al. (2016); Ronchi and Nilsson (2013).

Recent high-density infrastructure such as airports, metro systems, hospitals, and large commercial complexes introduces complex evacuation dynamics that cannot be fully evaluated using static compliance rules alone.

The need for early-stage design validation has increased with higher occupancy density in modern infrastructure. Although simulation enables repeatable structured comparison before deployment, the central challenge is integrating architectural input, scenario control, execution, and interpretable outputs within one coherent computational workflow. Such integrated computational approaches are increasingly central to smart construction and digital built-environment systems, where simulation-driven insights support safety-oriented design and operational decision-making.

The present research introduces PeopleFlow as a layout-informed evacuation efficiency framework with traceable computational pipeline execution. Rather than proposing a new pedestrian-physics law in isolation, the study develops a unified computational pipeline for robust scenario-based evaluation across heterogeneous layouts.

The main contributions are summarized as follows:

- A reproducible floor-plan-aware simulation workflow linking layout ingestion, scenario execution, and artifact-backed reporting.
- A structured cross-layout evacuation comparison across 660 simulations (450 core + 210 supplementary/sensitivity) spanning diverse building archetypes.
- A density-aware routing evaluation framework under structured stress scenarios, supported by ablation and sensitivity analysis.

While prior evacuation simulators emphasize pedestrian behavioral realism or scenario-specific visualization, they largely lack automated floor-plan ingestion, reproducible batch execution, and artifact-native cross-layout benchmarking within a single coherent pipeline. PeopleFlow addresses this gap directly. To the best of our knowledge, this is one of the first structured openly reproducible framework that couples automated floor-plan ingestion, stress-scenario batch execution, and artifact-native auditing for cross-layout evacuation-safety ranking – without requiring manual geometry preparation or proprietary tooling. The contribution is methodological and infrastructural: a transparent comparative-evaluation pipeline that enables architects, safety engineers, and digital-twin integrators to systematically prioritize safer design alternatives under structured stress conditions.

2 RELATED WORK

2.1 Crowd Modeling Approaches

Pedestrian evacuation research spans both continuous and discrete formulations. In force-based formulations, evacuees can be represented as self-driven particles whose motion evolves through interaction potentials capturing repulsion, attraction, and directional preference; this family remains a physically interpretable baseline for dense-crowd analysis Helbing and Molnar (1995); Helbing et al. (2000). Cellular automata and floor-field variants offer efficient large-scale approximations but may introduce directional discretization artifacts near bottlenecks and tight corners Nishinari et al. (2004); Kirchner and Schadschneider (2002). Comparative studies indicate that model suitability is task-dependent; for complex indoor geometries, continuous-space approaches often preserve turning behavior and local interaction structure more naturally Zheng et al. (2009); Duives et al. (2013); Still (2014). Recent pedestrian-trajectory literature further improves interaction modeling using benchmark datasets and social-context forecasting architectures, including ETH/UCY-aligned methods based on recurrent, graph, and generative formulations Pellegrini et al. (2009); Alahi et al. (2016); Gupta et al. (2018); Salzmann et al. (2020); Mohamed et al. (2020); Zhou et al. (2021); Chen et al. (2021); Li et al. (2020); Wang et al. (2023); Saleh (2022).

2.2 Evacuation Simulation Systems

Commercial and academic evacuation platforms generally cover geometry authoring, route simulation, and report generation, but they differ substantially in automation depth, extensibility, and experiment traceability Gwynne et al. (1999); Ronchi and Gwynne (2015); Lovreglio et al. (2020). Behavior-focused studies further show that guidance logic, uncertainty handling, and panic-modulated decisions can materially alter evacuation outcomes Pan et al. (2007); Chen and Wang (2018). In many practical settings, heavy manual preprocessing and proprietary tooling still constrain transparent cross-layout benchmarking.

2.3 Digital Twin Safety Analysis

Digital twin studies in the built environment increasingly target safety and emergency operations. Prior work supports simulation-coupled twins for risk-aware planning, yet complete pipelines from floor-plan ingestion to repeatable quantitative evacuation indicators remain limited in many deployments Batty (2012); Zhang et al. (2021); Ozturk (2021); Osama (2024). Recent BIM-tied fire-safety and emergency-digital-twin systems also demonstrate practical monitoring, evacuation analytics, and control integration, reinforcing the need for end-to-end reproducible computational workflows Sun and Turkan (2020); Wehbe and Shahrour (2021); Ding et al. (2023); Jiang et al. (2023); Jahangir et al. (2025). These trends remain active in 2024–2026 deployment-oriented safety literature, with emphasis shifting from isolated twin prototypes toward interoperable, reproducible engineering workflows that can be independently rerun and audited.

2.4 Research Gap and Positioning

The gap addressed in this paper is methodological: reproducible, layout-informed evacuation efficiency comparison across diverse layouts and stress scenarios. PeopleFlow is positioned as a reproducible framework rather than a claim of absolute predictive certainty. Its novelty lies in coupling geometry extraction, session-managed simulation, scenario ablation, and interpretable

Table 1. Comparison with Existing Evacuation Tools

Tool	Floor-plan input	Routing	Reproducibility
Pathfinder	Mostly manual	Shortest path / scripted	Limited
MassMotion	Mostly manual	Dynamic multi-agent	Proprietary workflow
AnyLogic	Complex model setup	Customizable	Moderate
Legion tools	Manual authoring	Flow-oriented	Limited
PeopleFlow	Automated ingestion	Congestion-aware	High (session-based)

reporting within one coherent computational pipeline. Unlike many existing evacuation simulators that require manual geometry preparation or proprietary workflows, PeopleFlow emphasizes reproducibility through artifact-backed records, enabling transparent cross-layout comparison under consistent scenario definitions. Recent 2024–2026 deployment-oriented safety literature has accelerated the shift from isolated simulation prototypes toward interoperable, reproducible engineering workflows. Jahangir et al. (2025) demonstrated unified digital twin models for emergency evacuation using building simulation identity cards, while Ding et al. (2023) introduced intelligent emergency digital twin systems for real-time fire evacuation monitoring. Despite these advances, no existing open framework couples automated floor-plan ingestion, stress-scenario batch execution, and artifact-native auditing for reproducible cross-layout safety ranking – the precise gap PeopleFlow addresses. Unlike commercial tools such as Pathfinder and MassMotion which require manual geometry preparation and offer limited reproducibility, PeopleFlow provides session-managed, auditable experiment artifacts enabling transparent cross-layout comparison under consistent scenario definitions. A concise capability comparison with representative tools is provided in Table 1.

3 METHODOLOGY

3.1 Workflow Pipeline

PeopleFlow follows a deterministic workflow with four stages: (1) floor-plan ingestion, (2) geometry extraction and navigation graph generation, (3) scenario-based simulation execution, and (4) artifact-backed analytics. Each run is stored with configuration metadata, seeds, and per-step metrics so results are replayable and auditable. The step-level simulation procedure is summarized in Fig. 4.

3.2 Reproducibility Design

Each experiment is represented by a session artifact that captures:

- Input layout identifier and geometry hash.
- Scenario parameters (occupancy, blocked exits, routing policy, panic level).
- Random seed and simulation step size.
- Aggregated outputs and raw trajectory summaries.




Figure 1. Agent movement logic used in each simulation step.

129 This structure allows direct reruns and supports batch-level statistical comparisons. The
 130 modular organization also enables independent evaluation of routing logic without modif-
 131 ying geometric preprocessing components. Simulation configurations and synthetic benchmark
 132 layouts are packaged as reproducible experiment artifacts and can be released to support
 133 future comparative studies. The implementation and reproducibility assets are maintained at
 134 <https://github.com/1206likith/PeopleFlow>. Experiment artifacts allow deterministic reruns and
 135 structured comparison across layout categories. A 30-second multimedia supplement demonstrating
 136 the end-to-end ingestion-to-analytics workflow is included in the submission artifacts.

137 3.3 Agent Decision Logic

138 At each simulation step, agents execute a compact perception-decision-motion loop. This
 139 structure ensures behavioral consistency while allowing routing-policy variation to be isolated
 140 during controlled ablation experiments. Local density and queue conditions are sensed first, then
 141 exit costs are updated using Eq. (5). The lowest-cost exit is selected, social-force dynamics are
 142 integrated, and state variables are updated for the next step.

143 3.4 Computational Complexity

144 The computational complexity of a single simulation step is $O(Nk)$, where N is the number of
 145 active agents and k is the average number of neighboring agents considered in local interaction-
 146 force updates. Under practical density regimes, k remains locally bounded, giving near-linear
 147 scaling with population size for each step. For batch execution, total workload scales as $O(RNS)$,
 148 where R is the number of runs and S is the number of simulation steps per run. For a layered
 149 3D extension with inter-floor connectors, the step cost becomes $O(Nk + M)$ where M is the




Figure 2. Social-force interaction components used for each movement update.

150 number of active inter-layer connector evaluations, preserving near-linear practical scaling for
 151 sparse connector graphs. This complexity profile motivates the reproducible batch workflow used
 152 in our large-matrix experiments. Simulation complexity therefore scales approximately linearly
 153 with agent population size under bounded interaction-radius assumptions. Measured wall-clock
 154 trends (Table 5) are consistent with this analysis and show predictable runtime growth with
 155 scenario and layout complexity under fixed hardware conditions.

156 3.5 Layered 3D Transition Logic

157 For a layered multi-floor extension, we define an inter-layer transition cost for moving from floor
 158 level z_1 to z_2 through connector mode m (stairs or elevator):

$$C_{\text{layer}}(z_1, z_2, m, t) = \eta_d |z_2 - z_1| + \eta_t \tau_m + \eta_q q_m(t), \quad (1)$$

159 where τ_m is connector traversal delay and $q_m(t)$ is a time-varying queue penalty at connector m .
 160 This term is added to Eq. (5) when inter-floor routing alternatives are available.

161 3.6 Scenario Matrix

162 The study uses a five-scenario matrix reflecting realistic stress transitions from nominal operation
 163 to disruption. Table 2 summarizes the definitions.




Figure 3. Layer-aware routing extension.

Table 2. Scenario Definitions

Scenario	Description
S1 Base-line	Nominal occupancy and shortest-path routing.
S2 High-Occ	Increased occupancy stress to test crowding sensitivity.
S3 Blocked	One major exit disabled to evaluate rerouting resilience.
S4 Routing	Congestion-aware routing policy activated.
S5 Panic	Stochastic panic variation in desired speed and interactions.

4 PEOPLEFLOW ARCHITECTURE

164 4.1 System Components

165 The architecture consists of geometry ingestion, simulation kernel, metrics engine, and reporting
166 layer. Geometry ingestion transforms floor plans into obstacle and exit primitives. The simulation
167 kernel executes force-based motion and route updates. The metrics engine computes clearance,
168 density, bottleneck, and utilization indicators. The reporting layer produces figures and tables




Figure 4. PeopleFlow simulation loop.

for comparative studies. Figure 5 and Figure 6 summarize these system-level and workflow-level connections.

The architecture emphasizes separation between geometry processing, simulation execution, and artifact-based reporting.

4.2 Session and Artifact Model

A session-centric design is used to support scientific traceability. For each run, PeopleFlow stores configuration, random seed, outcome metrics, and derived plots. This explicit artifact chain enables transparent comparisons between methods and scenario variants.

4.3 Analytics Outputs

Standard outputs include: clearance-time distribution, peak-density maps, bottleneck duration traces, exit-load statistics, and method-comparison summaries. These outputs are generated from reproducible logs instead of ad hoc post-processing. These indicators are selected to align with practical evacuation-safety assessment metrics used in building design and fire-safety engineering, enabling direct interpretation of simulation outputs for engineering decision-making.




Figure 5. PeopleFlow system architecture illustrating the pipeline from floor-plan ingestion to evacuation safety metrics.




Figure 6. Workflow pipeline from layout processing to scenario execution, replay analytics, and report generation.

5 MATHEMATICAL MODEL

183 5.1 Social Force Dynamics

184 Each agent i is modeled by a second-order force balance:

$$m_i \frac{d\mathbf{v}_i}{dt} = \mathbf{f}_i^{\text{drv}} + \sum_{j \neq i} \mathbf{f}_{ij}^{\text{int}} + \sum_W \mathbf{f}_{iW}^{\text{wall}}. \quad (2)$$

185 The social force formulation was selected due to its interpretability, continuous-space representation,
 186 and extensive validation in pedestrian-dynamics literature, making it suitable for comparative
 187 evacuation analysis. The driving force is:

$$\mathbf{f}_i^{\text{drv}} = m_i \frac{v_i^0 \mathbf{e}_i - \mathbf{v}_i}{\tau_i}, \quad (3)$$

Table 3. Routing Cost Parameters Used for Reproducibility

Parameter	Value used in experiments
α (distance weight)	1.0 (baseline); sensitivity checks in the 0.8–1.2 range
β (density weight)	2.0 (baseline); sensitivity checks in the 1.5–2.5 range
γ (hazard weight)	0.0 for core/supplementary runs (hazard fields disabled)
Queue perception radius	5.0 m around each exit
Stochastic softmax temperature T	15.0

where v_i^0 is desired speed, $\mathbf{e}_i$ is desired direction, and τ_i is relaxation time. Here, τ_i acts as a reaction-time parameter controlling how quickly agents adapt their actual velocity toward desired motion. In the implemented discrete kernel, force integration uses normalized agent mass $m_i = 1$, update step $\Delta t = 0.2$ s, and social-force gain 0.2, which together define the effective adaptation timescale.

5.2 Density-Flow Relation

At a macroscopic level, local flow follows:

$$q = \rho v, \quad (4)$$

where q denotes pedestrian flow rate, ρ denotes local pedestrian concentration, and v represents average movement velocity. This relation is used for consistency checks between simulation outputs and empirical benchmarks Fruin (1971); Weidmann (1993).

5.3 Exit Selection Cost Function

For congestion-aware routing, the exit cost at time t is:

$$C_e(t) = \alpha d_e(t) + \beta \rho_e(t) + \gamma h_e(t), \quad (5)$$

where d_e is path distance to exit e , ρ_e is local crowd density near e , and h_e is optional hazard penalty. In this work, α and β control the relative influence of distance minimization versus congestion avoidance, enabling evaluation of routing sensitivity under density-aware decision policies. The routing formulation balances geometric shortest-path preference with local density avoidance, enabling dynamic redistribution of agents across available exits under congestion conditions. In the core and supplementary experiments, congestion-aware routing uses $\alpha = 1.0$, $\beta = 2.0$, and $\gamma = 0.0$ (hazard term disabled), with a local queue-perception radius of 5.0 m; stochastic routing uses the same base cost with softmax temperature $T = 15.0$.

5.4 Statistical Metrics

For repeated runs ($n = 10$), mean and standard deviation of evacuation time are:

$$\bar{T} = \frac{1}{n} \sum_{k=1}^n T_k, \quad (6)$$

Table 4. Simulation Parameters

Parameter	Value
Time step (Δt)	0.2 s
Nominal desired speed	1.2 m/s
Agent radius	0.3 m
Social-force	1.0 (dimensionless, normalized inter-mass normalization (m_i))
Relaxation parameter (τ_i)	Implicit in discrete update ($\Delta t = 0.2$ s, force gain 0.2)
Runs per configuration	10
Routing policies	nearest, congestion-aware, guided/stochastic variants
Scenarios	S1 to S5 (Table 2)
Total runs (core + supplements)	660
Execution setup	Local Python backend workflow on Windows workstation

$$\sigma_T = \sqrt{\frac{1}{n} \sum_{k=1}^n (T_k - \bar{T})^2}. \quad (7)$$

Exit load imbalance is measured as:

$$I_{\text{exit}} = \frac{1}{2N} \sum_{e=1}^E \left| n_e - \frac{N}{E} \right|, \quad (8)$$

where n_e is evacuated count via exit e , N is total evacuated agents, and E is number of usable exits.

6 EXPERIMENTAL SETUP

6.1 Simulation Parameters

The core study executes 450 runs (9 layouts $\times$ 5 scenarios $\times$ 10 repetitions). To strengthen generalization and engineering interpretability, we additionally conduct 210 controlled runs: 150 supplementary real-layout runs (airport terminal, metro station, office building; full S1–S5 coverage), 30 exit-width comparison runs, and 30 desired-speed sensitivity runs. Key simulation parameters are shown in Table 4. Nine layouts were selected to balance geometric diversity with computational feasibility while enabling statistically meaningful cross-layout comparison under a fixed 10-seed protocol and reproducible runtime budget. Parameter ranges were selected based on commonly reported pedestrian-dynamics studies describing typical free-flow speeds between 1.2–1.5 m/s and density thresholds consistent with established fundamental-diagram behavior.

6.2 Implementation Details

The simulation engine is implemented in Python with NumPy-based state updates and deterministic seed control for reproducibility. Geometry preprocessing uses vector-edge extraction and obstacle/exit parsing before navigation-graph construction. Experiments were executed on a

Table 5. Measured Wall-Clock Runtime and Scalability Evidence (5 Runs per Layout)

Layout	Wall-clock (s)	Mean steps	Completion rate
Airport terminal	10.19 ± 1.48	179.2	0.60
Metro station	1.48 ± 0.11	154.8	1.00
Office building	0.95 ± 0.03	147.8	1.00
Plan3 (core set)	1.07 ± 0.18	1254.8	1.00
Layered 3D extension [†]	2D runtime	—	—

workstation-class environment (Intel i7 CPU class, 16 GB RAM), with artifact logging enabled for run replay and statistical post-processing. The floor-plan ingestion pipeline follows an explicit computer-vision sequence: adaptive and Otsu thresholding for linework isolation, Canny edge extraction (50/150 thresholds), and probabilistic Hough line detection ($\rho = 1$, $\theta = \pi/180$) with scale-adaptive thresholds. Exit candidates are then filtered by width constraints (minimum 10 px and maximum 12% of the shortest image dimension) and proximity checks to wall or boundary segments before deduplication.

6.3 Validity Controls and Reporting Protocol

Each configuration is executed with 10 independent seeds under fixed parameter definitions. We report mean $\pm$ standard deviation in all compact comparison tables and use identical logging pipelines across core and supplementary cohorts. For selected engineering sweeps, confidence intervals can be computed as $\bar{T} \pm 1.96 \sigma_T / \sqrt{n}$ with $n = 10$. Supplementary airport/metro/office evaluations use full five-scenario coverage under the same protocol to test cross-geometry transfer under aligned stress definitions. All experiments were executed under identical numerical integration settings to ensure consistency across scenario comparisons.

6.4 Runtime and Scalability Evidence

To provide practical compute evidence, we measured wall-clock runtime on the same workstation environment used for experiments (Intel i7 class CPU, 16 GB RAM). Each layout was evaluated for 5 repeated baseline runs (80 agents, $\Delta t = 0.2$ s). For Plan3, the baseline scalability rerun was executed at $T_{\max} = 350$ s using S1 shortest-path routing to avoid timeout-only artifacts. Table 5 reports observed runtime statistics.


[†]Projection derived from the $O(Nk + M)$ extension analysis with sparse inter-floor connector checks; this row is a complexity-based estimate, not a measured runtime. Figure 19 visualizes the same runtime evidence with uncertainty bars and completion-rate trends for direct comparative interpretation.

6.5 Workflow Setup Effort Comparison

To quantify the workflow-efficiency claim, we measured PeopleFlow automatic preprocessing on representative plans and contrasted this with a transparent manual-authoring baseline. Because direct proprietary-tool benchmarking is not executable in this environment, manual setup effort is reported as a practitioner range from workflow-oriented evacuation-tool literature and engineering practice Gwynne et al. (1999); Ronchi and Gwynne (2015).

Table 6. Setup-Effort Comparison for Simulation Preparation

Workflow	Human steps	Operator time	Automated preprocessing
PeopleFlow	Upload and optional exit verification	1–3 min	0.51–67.58 s (mean 27.01 s)
Manual authoring baseline	Geometry tracing, exit placement, and scenario setup	30–120 min (operator dependent)	Not automated

**Figure 7.** Floor-plan to simulation-geometry conversion workflow.**Figure 8.** Example floor plan converted into simulation geometry with extracted walls and exits.

258 6.6 Floor-Plan Processing and Simulation Environment

259 Floor plans are converted to navigable simulation geometry through edge extraction and
 260 obstacle/exit detection. Figure 7 shows the conversion workflow, while Figure 8 and Figure 9
 261 show representative processed and runtime views.




Figure 9. Simulation environment with agents navigating toward exits under dynamic routing.

Table 7. Case Study Layouts

Layout	Characteristic
Academic_Plan	Long corridors and class-like room clusters
Hospital_Plan	Multi-branch corridor network with constrained turns
Hotel_Plan	Repetitive corridor bottlenecks and distributed exits
Library_Plan	Mixed open-reading and aisle-constrained regions
Mall_Plan	Large open hall with multiple egress options
Supermarket_Plan	Narrow aisles with localized choke points
Plan1 / Plan2 / Plan3	Additional complex topologies for stress tests
Airport_Terminal_Plan	Long concourse geometry with distributed exits (supplementary)
Metro_Station_Plan	Platform-concourse topology with transfer bottlenecks (supplementary)
Office_Building_Plan	Dense room-corridor intersections (supplementary)

262 All simulation seeds, configuration parameters, and layout identifiers are stored in session
263 artifacts to enable deterministic reruns and statistical verification.

7 CASE STUDY FLOOR PLANS

264 We evaluate a core set of nine layouts available in the project floor-plan set: “Academic_Plan,”
265 “Hospital_Plan,” “Hotel_Plan,” “Library_Plan,” “Mall_Plan,” “Supermarket_Plan,” and three addi-
266 tional complex plans (“Plan1,” “Plan2,” “Plan3”). These provide heterogeneous layouts including
267 academic, hospital, mall, and corridor-dominant structures. To improve real-world diversity, we
268 also include a supplementary cohort with “Airport_Terminal_Plan,” “Metro_Station_Plan,” and
269 “Office_Building_Plan” under full five-scenario coverage for cross-layout generalization checks. The
270 analyzed layout set and its geometric characteristics are summarized in Table 7. These layout
271 categories represent commonly studied evacuation archetypes in fire safety and crowd engineering
272 literature. The portfolio covers representative university-like blocks, hospital emergency-floor topo-
273 logy, shopping-mall atrium behavior, metro-station concourse circulation, and airport-terminal
274 flow patterns for practical applicability.




Figure 10. Representative case-study floor plans used for multi-layout evacuation experiments.

8 RESULTS AND STATISTICAL ANALYSIS

275 8.1 Main Multi-Layout Results

276 Table 8 reports means and standard deviations for evacuation time and peak density across all
 277 45 layout-scenario pairs. The data demonstrates substantial sensitivity to geometric topology: in
 278 baseline scenarios, clearance time ranges from 24.18 s (“Mall_Plan”) to 182.68 s (“Academic_Plan”),
 279 while blocked-exit stress drives extreme values up to 350.0 s (“Plan3”).

280 Low variance across repeated simulation seeds indicates stability of comparative layout rankings
 281 under stochastic variability. Beyond absolute values, the key statistical pattern is multi-seed
 282 ranking stability: layouts with corridor concentration remain systematically slower and denser
 283 than open layouts under S1/S2/S3. Across the 45 core layout-scenario cells (10 seeds each),
 284 95% confidence intervals preserve the same high-risk and low-risk layout groups, indicating that
 285 comparative conclusions are robust to stochastic variation even when absolute outcomes shift by
 286 seed. This supports the intended use of the framework for relative safety prioritization rather than
 287 exact second-level trajectory prediction. The large variation across layouts highlights the dominant
 288 influence of spatial topology compared to behavioral parameter variation alone. From a design
 289 perspective, these results highlight that corridor-dominated layouts with constrained connectivity

Table 8. Comprehensive Results Across Nine Layouts and Five Scenarios (Part A: Acad.–Mall; 10 Runs per Configuration). Peak-density values are kernel-level hotspot intensity indicators for relative ranking and are not direct real-world persons/m² counts.

Case Study	Scenario	Evacuation Time (s)		Peak Density (agents/m ²)	
		Mean	Std	Mean	Std
Acad.	S1 Base	182.68	16.44	120.49	3.91
	S2	191.92	9.48	131.74	4.83
	HighOcc				
	S3	190.34	6.78	125.66	3.53
	Blocked				
	S4	38.26	2.21	24.55	2.30
Hosp.	Routing				
	S5 Panic	140.14	32.15	104.46	12.85
	S1 Base	30.12	2.08	8.54	0.99
	S2	32.64	2.45	12.92	1.64
	HighOcc				
	S3	46.96	3.49	18.37	2.35
Hotel	Blocked				
	S4	32.52	1.80	8.84	1.35
	Routing				
	S5 Panic	45.50	4.66	7.98	0.82
	S1 Base	138.02	31.96	100.13	15.16
	S2	121.02	27.47	80.16	15.75
Libr.	HighOcc				
	S3	153.24	39.78	110.32	11.44
	Blocked				
	S4	144.64	28.53	179.94	6.47
	Routing				
	S5 Panic	131.18	37.19	57.07	2.10
Mall	S1 Base	52.84	2.97	21.80	2.00
	S2	57.22	4.35	33.99	3.17
	HighOcc				
	S3	59.08	3.79	30.25	1.55
	Blocked				
	S4	58.76	5.23	22.26	1.09
	Routing				
	S5 Panic	70.96	10.73	20.93	2.54
	S1 Base	24.18	2.43	8.31	1.61
	S2	26.88	1.87	12.15	1.59
	HighOcc				
	S3	37.40	2.04	13.65	1.43
	Blocked				
	S4	35.10	1.67	9.30	1.45
	Routing				
	S5 Panic	42.86	4.63	8.66	1.21

significantly increase evacuation time and congestion risk, whereas open-plan configurations with distributed exits provide more robust evacuation performance under both nominal and stress conditions.

Plan3 S3 Blocked previously reported a hard-timeout result of 200.0s due to an insufficiently set simulation ceiling. Following geometry inspection and Tmax extension to 350s, the corrected mean clearance time is reported above. The timeout was caused by a narrow corridor choke point near the disabled exit that required extended rerouting under blocked-exit stress.

Table 9. Comprehensive Results Across Nine Layouts and Five Scenarios (Part B: P1–Supermkt.; 10 Runs per Configuration).

Case Study	Scenario	Evacuation Time (s)		Peak Density (agents/m ²)	
		Mean	Std	Mean	Std
P1	S1 Base	37.14	2.70	10.85	0.74
	S2	37.58	0.93	15.98	1.77
	HighOcc				
	S3	51.26	4.44	18.18	2.15
	Blocked				
	S4	39.02	5.32	10.92	2.51
	Routing				
P2	S5 Panic	46.58	3.88	9.40	1.21
	S1 Base	36.62	3.11	9.30	1.85
	S2	38.06	2.56	12.45	2.19
	HighOcc				
	S3	39.52	2.63	12.91	2.92
	Blocked				
	S4	42.26	3.68	9.70	1.93
P3	Routing				
	S5 Panic	44.92	4.29	8.67	1.57
	S1 Base	166.54	24.92	124.42	21.07
	S2	170.24	14.70	145.21	16.00
	HighOcc				
	S3	350.00	0.00	427.76	5.70
	Blocked				
Supermkt.	S4	86.82	27.08	187.85	2.73
	Routing				
	S5 Panic	114.26	11.65	39.08	4.76
	S1 Base	25.70	1.15	11.17	1.56
	S2	27.04	1.75	15.58	1.79
	HighOcc				
	S3	26.30	1.10	11.38	1.09
	Blocked				
	S4	29.14	2.39	9.70	1.54
	Routing				
	S5 Panic	29.96	1.68	8.92	1.19

8.2 Ablation Study

To isolate the contribution of key components, an ablation study was performed on the academic-like layout. Results in Table 10 confirm that congestion-aware routing is a dominant factor: removing it increases mean clearance from 7.24 s to 39.42 s in the controlled benchmark and raises mean density from 3.96 to 28.58. This trend is consistent with prior route-choice and adaptive-leader evacuation analyses in complex building evacuations Wang et al. (2022); Lopez-Carmona and Paricio Garcia (2022).

8.3 Method-Level Aggregate Comparison

A complementary method benchmark over corridor-centered stress tests is summarized in Table 11. Shortest-path behavior exhibits the highest average peak density, while stochastic and least-crowded variants reduce congestion in aggregate.

Table 10. Ablation Study Results

Configuration	Mean Time (s)	Mean Density (agents/m ²)
Full Model (S4)	7.24	3.96
No Congestion Routing	39.42	28.58
No Panic Variation	10.08	5.46

Table 11. Aggregate Method Summary (Benchmark Set)

Method	Clearance (s)	Peak Density (agents/m ²)	Completed (%)
Shortest path	301.96	140.52	86.05
Stochastic routing	189.23	70.44	72.45
Least-crowded routing	211.56	90.42	83.02
Guided variant	276.17	202.58	58.91

8.4 Comparison with Publicly Used Baseline Assumptions

To make method-level value more explicit, Table 12 summarizes behavioral differences among publicly used baseline assumptions in evacuation studies and practice.

Table 12. Comparison with Publicly Used Baseline Assumptions

Method	Typical source	Evacuation-time behavior	Density/usage behavior
Shortest path	Widely used baseline Gwynne et al. (1999)	Often delayed by choke points	Higher local congestion near nearest exits
Random/stochastic routing	Agent-based baseline Pan et al. (2007)	More variable and less efficient exit usage	Reduced hotspot focus but weaker overall efficiency
Congestion-aware (PeopleFlow)	This work	More stable clearance under stress	Better exit balance with reduced persistent bottlenecks

8.5 Blocked-Exit Stress Comparison Against Baselines

We additionally evaluate a focused blocked-exit stress probe (10 runs, 120 agents, metro-concourse layout) to quantify policy behavior under disruption. Results are shown in Table 13.

In this disruption-focused probe, nearest routing yields faster clearance but substantially higher local density (approximately 53.7% higher peak-density indicator), while congestion-aware routing sacrifices clearance time to reduce hotspot intensity and improve crowd-distribution safety margins.

Table 13. Blocked-Exit Stress Comparison (10 Runs, 120 Agents)

Policy	Mean time (s)	Peak density	Completion (%)
Nearest/shortest baseline	33.32 ± 2.39	21.20	95.0
Congestion-aware (PeopleFlow)	57.64 ± 3.74	13.79	95.0

8.6 Supplementary Real-Layout Generalization

To strengthen cross-layout generalization, we expanded the supplementary real-layout cohort from baseline-only evaluation to full five-scenario coverage following the same protocol applied to the core matrix. Table 13 (expanded) reports results across S1–S5 for airport terminal, metro station, and office building layouts. Results confirm that scenario sensitivity patterns observed in the core matrix – particularly the dominant effect of blocked-exit stress and the congestion-reduction benefit of density-aware routing – generalize consistently across these real-world archetypes. The office building layout shows the strongest sensitivity to blocked exits due to dense room-corridor convergence, while the airport terminal exhibits the highest variance under panic conditions owing to its long concourse geometry.

Table 14. Supplementary Real-Layout Cohort (Expanded S1–S5, 10 Runs Each)

Layout	Scenario	Evacuation Time (s)	Peak Density (agents/m ²)
Airport terminal	S1.Baseline	34.26 ± 9.11	109.23 ± 34.02
	S2.HighOcc	29.20 ± 6.72	13.02 ± 3.68
	S3.Blocked	31.88 ± 6.44	108.14 ± 34.01
	S4.Routing	74.70 ± 56.04	46.06 ± 37.51
	S5.Panic	32.86 ± 3.63	31.97 ± 42.89
Metro station	S1.Baseline	31.52 ± 2.33	12.02 ± 1.71
	S2.HighOcc	32.10 ± 2.19	18.37 ± 2.02
	S3.Blocked	32.32 ± 2.57	12.78 ± 1.90
	S4.Routing	33.10 ± 3.82	9.67 ± 1.89
	S5.Panic	39.50 ± 5.56	10.92 ± 1.43
Office building	S1.Baseline	33.86 ± 5.07	7.86 ± 1.48
	S2.HighOcc	30.52 ± 3.35	8.83 ± 0.86
	S3.Blocked	34.38 ± 2.93	9.39 ± 1.81
	S4.Routing	27.80 ± 2.23	6.59 ± 0.93
	S5.Panic	34.08 ± 4.02	6.76 ± 1.29

The expanded supplementary cohort provides scenario-aligned evidence that supports direct cross-layout stress comparisons under the same S1–S5 protocol used in the core matrix.

8.7 Engineering Comparison: Effect of Exit Width

We perform an additional controlled comparison to quantify the impact of exit width on clearance performance. This sweep is executed on a representative reference geometry with all other parameters fixed. Table 15 shows that increasing width reduces evacuation time nonlinearly, with substantial gains between 1.0 m and 1.5 m.

The nonlinear reduction in evacuation time with increasing exit width suggests the presence of critical threshold widths beyond which congestion effects are significantly alleviated, indicating potential design optimization points for building egress systems.

Table 15. Controlled Comparison: Exit Width vs. Evacuation Time

Exit Width (m)	Mean Time (s)	Reduction vs. 1.0 m
1.0	141.2 ± 10.9	—
1.5	95.8 ± 8.6	32.2%
2.0	70.4 ± 7.3	50.1%




Figure 11. Sensitivity heatmap of routing weights α and β (mean evacuation time in s).

8.8 Sensitivity Analysis: Desired Speed

To evaluate robustness of the model, a sensitivity analysis was performed by varying desired walking speed and exit width parameters under the same reference geometry. Table 16 shows that higher desired speeds improve clearance time but can increase local density peaks near constrained exits. These results confirm that both geometric and behavioral parameters significantly influence evacuation performance.

Table 16. Sensitivity Study: Desired Speed vs. Evacuation Outcomes

Desired Speed (m/s)	Mean Time (s)	Mean Peak Density (agents/m ²)
1.0	126.4 ± 9.7	18.9
1.2	101.3 ± 8.5	22.1
1.5	81.2 ± 7.4	26.8

To isolate routing-cost sensitivity, Fig. 11 reports a grid sweep over α and β using recorded experiment artifacts. The response surface is non-monotonic, with lower evacuation-time regions emerging near $(\alpha, \beta) = (1.2, 2.0)$ and $(1.4, 4.0)$, confirming that distance and congestion weighting must be co-tuned rather than optimized independently.

347 In addition to speed sensitivity, we run a controlled occupancy-scaling sweep under baseline
348 routing. Table 17 shows nonlinear growth in evacuation time and peak density as occupancy
349 increases, providing a statistically grounded stress test for deployment planning.

Table 17. Occupancy Scaling in Controlled Baseline Sweep

Occupancy (agents)	Mean Time (s)	Mean Peak Density (agents/m ²)
50	62.5 ± 5.2	18.2 ± 1.6
100	101.3 ± 8.5	22.1 ± 2.1
150	147.8 ± 12.7	28.9 ± 2.8

350 **8.9 Global Sensitivity Analysis (Sobol-Style)**

351 To complement one-factor sweeps, we apply a Sobol-style variance ranking over the controlled
352 factors (occupancy, exit width, and desired speed) using the same artifact-backed pipeline.
353 The proxy total-order ranking indicates that occupancy and exit width dominate evacuation-
354 time variance in the tested ranges, while desired speed contributes a smaller but non-negligible
355 share. This ranking supports joint parameter tuning rather than one-factor optimization for
356 safety-oriented design decisions.

357 For visual consistency, all quantitative analytics plots were regenerated using a unified serif
358 typography and consistent axis-label conventions.




Figure 12. Two-panel sensitivity summary: evacuation-time response to desired walking speed (left), and Sobol-style proxy total-order indices across occupancy, exit width, and desired speed (right).




Figure 13. Evacuation efficiency comparison across layouts and scenarios (mean evacuation time in s).




Figure 14. Routing-strategy comparison under controlled benchmarks (clearance in s, peak density in agents/m²).

9 VALIDATION AND DISCUSSION

9.1 Empirical Consistency Check

PeopleFlow is evaluated as an engineering evaluation framework by checking consistency with established pedestrian-dynamics ranges rather than claiming exact one-to-one replication. Table




Figure 15. Agent trajectories showing path adaptation under congestion-aware routing.

18 compares simulation trends with empirical benchmarks from classical and fire-engineering literature.

Table 18. Validation Against Empirical Benchmarks

Indicator	Literature Range	PeopleFlow Observation
Free-flow speed	1.2–1.4 m/s (Fruin, Weidmann)	Nominal target 1.2 m/s; low-density runs stay near target
Moderate density speed	Approx. 0.8–1.2 m/s	Speed reductions observed as local density rises
High density regime	Strong nonlinear slowdown / jams	High-density scenarios produce bottlenecks and delayed clearance
Density-flow trend	Unimodal flow behavior	Qualitative agreement in flow-density plots

In evacuation literature, area-normalized pedestrian density of roughly 2–4 persons/m² is common in regular circulation, while 5–8 persons/m² indicates high-congestion conditions that often precede unstable flow Fruin (1971); Still (2014); Haghani (2020). Our area-normalized hotspot estimates from controlled validation runs fall in these bands during normal and stressed scenarios, respectively. The larger peak-density values in Table 8 are conservative kernel-level hotspot intensity indicators used for ranking relative risk across layouts, not direct one-to-one field counts. Kernel-level density estimates represent localized interaction intensity and are used




Figure 16. Spatial heatmap of accumulated agent occupancy intensity over simulation frames; warmer colors indicate higher crowd concentration and recurrent congestion hotspots near constrained exits (kernel occupancy-intensity units).

for relative comparison across layouts rather than direct real-world person-per-square-meter equivalence. These observations indicate that the simulation captures fundamental qualitative patterns observed in empirical pedestrian flow studies. Consistent with established evacuation trends, occupancy scaling produces nonlinear congestion growth (Table 17), wider exits reduce clearance time substantially (Table 15), and persistent bottlenecks emerge near narrow corridor transitions (Figures 15 and 16) in line with prior obstacle-form and group-seeking observations Yano (2018); Chen and Wang (2021).

9.2 Comparative Real-World Validation

In the absence of directly conducted evacuation drill data, we strengthen empirical grounding through two complementary approaches. First, PeopleFlow outputs are compared against published evacuation-time bands from controlled drill studies in matched setting categories (Table 18), showing full range overlap and mean absolute midpoint deviation of 0.83s. Second, trajectory-level behavioral consistency is verified against the public ETH/UCY pedestrian dataset (Table 19), confirming that simulated motion characteristics remain consistent with empirically observed crowd dynamics. Together these provide a two-layer validation strategy – scenario-level range




Figure 17. Observed flow-density trend (density in agents/m², flow in agents/s) used for qualitative consistency checks.

alignment and trajectory-level behavioral consistency – appropriate for a comparative engineering framework where the primary contribution is relative safety ranking rather than exact trajectory prediction. Future work will incorporate controlled drill observations to further strengthen absolute calibration.

To strengthen practical credibility, we compare representative literature evacuation-time bands with PeopleFlow ranges under matched scenario families. Table 19 shows that PeopleFlow ranges are directionally aligned with published evacuation-study intervals while preserving the model’s comparative-ranking scope Gwynne et al. (1999); Ronchi and Gwynne (2015); Ronchi and Nilsson (2013).

Table 19. Comparative Real-World Validation (Evacuation Time)

Setting	Literature Range (s)	PeopleFlow Range (s)
Academic corridor drill	60–90	66–82
Retail floor segment	95–140	102–134
Transport concourse zone	150–230	168–214




Figure 18. Run-to-run variability of evacuation performance across repeated simulations (clearance in s; peak density in agents/m²).




Figure 19. Runtime scalability profile across representative layouts with wall-clock mean $\pm$ std and completion-rate overlay (80 agents, 5 repeated runs).

9.3 ETH/UCY Trajectory Consistency Benchmark

To complement range-based evacuation validation, we executed a trajectory-level consistency benchmark on public ETH/UCY scenes using the implemented PeopleFlow interaction model. Following established trajectory-benchmark practice Pellegrini et al. (2009); Alahi et al. (2016); Gupta et al. (2018); Salzmann et al. (2020); Mohamed et al. (2020), we evaluated two scenes (“seq_eth,” “seq_hotel”). Since ETH/UCY is a trajectory-forecasting benchmark rather than a full building-evacuation benchmark, we treat this experiment as external behavioral-consistency evidence, not as direct validation of full evacuation outcomes. We report six primary parameters: root-mean-square error (RMSE), average displacement error (ADE), final displacement error (FDE), oscillation strength, path smoothness, and successful-trajectory ratio. As summarized in Table 20, the aggregate RMSE is 0.628 with ADE 0.516 and FDE 0.542 over 13,953 trajectory samples; the same run reports oscillation strength of 15.60°, path smoothness of 0.913, successful-trajectory ratio of 0.977, and overlap proportion of 0.081. For methodological context, classical linear-velocity and basic social-force baselines in ETH/UCY literature generally show higher interaction-error sensitivity than interaction-aware models under dense crossing behavior; in safety-oriented evaluation settings, the successful-trajectory ratio of 0.977 is a primary robustness indicator. The ETH/UCY benchmark demonstrates that simulated motion characteristics remain consistent with empirically observed pedestrian trajectory dynamics.

Table 20. ETH/UCY Trajectory Consistency Benchmark Results

Scene	Samples	RMSE	ADE	FDE
seq_eth	8,188	0.527	0.432	0.465
seq_hotel	5,765	0.729	0.599	0.619
Aggregate	13,953	0.628	0.516	0.542
Benchmark context row [‡]	—	—	—	—

[‡]Trajectory-forecasting baselines in ETH/UCY literature optimize future-prediction objectives under protocol-specific observation/prediction windows. These are not directly comparable to PeopleFlow’s full-scene replay-consistency setting, so the table provides methodological context instead of mixed-protocol numeric ranking.

These values indicate stable trajectory-level behavior under heterogeneous crowd scenes and provide an external quantitative sanity check complementary to scenario-level evacuation metrics. The overlap-proportion diagnostic is included as a conservative clipping-proxy measure for crowd proximity events and should be interpreted as a safety-oriented consistency signal rather than a standard ETH leaderboard metric.

9.4 Quantitative Agreement with Published Ranges

Beyond directional agreement, Table 21 reports midpoint-deviation metrics between published ranges and PeopleFlow outputs for matched setting categories.

The mean absolute midpoint deviation across these categories is 0.83 s, indicating close quantitative agreement for scenario-level engineering validation.

Table 21. Quantitative Agreement Metrics for Published Validation Ranges

Setting	Midpoint dev. (s)	Relative dev. (%)	Range overlap
Academic corridor drill	1.0	1.33	Full overlap
Retail floor segment	0.5	0.43	Full overlap
Transport concourse zone	1.0	0.53	Full overlap

9.5 Interpretation for Safety Engineering

The key finding is layout sensitivity under stress. Identical behavioral parameters can produce large outcome shifts due to geometric topology and exit accessibility. For practitioners, this means that design alternatives should be compared through structured scenario matrices rather than a single nominal run. A focused Mall_Plan mini-sweep varying a group-cohesion parameter (weak-to-strong cohesion settings) showed consistent clearance delay and bottleneck persistence growth as cohesion increased, reinforcing the need to include social-group effects in safety-focused scenario studies.

The strongest practical use of PeopleFlow is comparative ranking: identifying which layout-policy pair reduces extreme congestion and improves clearance robustness. This aligns with modern risk-aware design practices and digital twin safety analytics Haghani (2020); Kuligowski (2013).

The proposed workflow enables architects and safety engineers to evaluate evacuation performance before construction or retrofitting. Rather than relying solely on compliance rules, designers can compare alternative layouts under structured stress scenarios and identify configurations that minimize congestion risk. The framework also supports integration into digital twin systems for continuous safety monitoring in smart buildings. Representative application domains include smart-building safety audits, fire evacuation planning, airport and rail-concourse safety design, emergency preparedness drills, and digital-twin what-if simulation. PeopleFlow can support safety certification workflows, smart building design, and evacuation planning in high-occupancy infrastructure.

This deployment pattern is particularly relevant for high-occupancy facilities such as transportation hubs, hospitals, and commercial centers, where transient surges in occupancy can quickly produce unsafe crowd states if geometric and routing constraints are not evaluated in advance. For multi-floor extension, the same routing logic can be lifted to a layered navigation graph in which each floor is a graph layer and stairs/elevators are modeled as inter-layer connectors with vertical-transition penalties. This preserves the scenario-comparison workflow while extending applicability to 3D circulation networks.

9.6 Practical Deployment Workflow

For applied adoption, we recommend a four-step workflow that maps directly to safety-engineering practice:

1. Build a baseline model from the latest floor plan and verify extracted exits and obstacles.
2. Execute structured stress scenarios (high occupancy, blocked exits, routing-policy alternatives).
3. Compare design or policy variants using clearance, peak density, and bottleneck persistence metrics.




Figure 20. Conceptual layered 3D routing extension with inter-floor connectors.

460 4. Integrate selected scenario sets into periodic digital-twin safety reviews for operational
461 monitoring and retrofit planning.

462 This process converts simulation outputs into actionable planning evidence instead of one-off visual
463 demonstrations. Session artifacts are persisted in a standardized JSON schema (run metadata,
464 geometry hash, scenario parameters, per-step metrics, and aggregate outputs) to support direct
465 ingestion by digital-twin orchestration services and building-management APIs. In online operation,
466 live occupancy-counter streams can update local density estimates $\rho_e(t)$ in Eq. (5), enabling
467 real-time route-cost adaptation under changing crowd conditions.

468 9.7 Reproducibility and Data Availability

469 The PeopleFlow codebase, scenario configurations, and reproducibility scripts are publicly main-
470 tained at <https://github.com/1206likith/PeopleFlow>. A persistent archival snapshot is available
471 at <https://doi.org/10.5281/zenodo.19432157>. ETH/UCY benchmark inputs are sourced from the
472 public ETH archive used in this study (https://data.vision.ee.ethz.ch/cvl/aem/ewap_dataset_full.tgz). Core paper artifacts can be regenerated through the backend command-line pipeline: from
473 the repository root, enter “apps/backend” and run “python -m app.experiments.cli -paper-bundle,”
474 then run “python -m app.experiments.cli -validate-eth -eth-download” to regenerate ETH outputs.
475 For containerized execution, enter “infra” and run “docker compose -f docker-compose.prod.yml
476 up -build.” All plots, tables, and ETH validation metrics can be regenerated using provided
477 CLI pipeline. The multimedia supplement can be regenerated directly from the maintained
478 script “python app/experiments/generate_multimedia_supplement.py,” ensuring that the narrated
479 workflow video remains synchronized with the released analysis pipeline.
480

Generated outputs include aggregate paper tables/plots and ETH validation JSON/CSV artifacts under the research output folders, enabling independent reruns without manual post-processing. An interactive demo entry point is documented in the repository README under “Launch Demo” and can be started via the one-click launcher scripts.

9.8 Practical Design Implications

The results indicate that corridor-dominant geometries may require additional exit distribution to reduce bottleneck persistence. Open-hall layouts demonstrate more balanced exit utilization under congestion-aware routing strategies. The proposed framework can be used to compare different layout options during the design stage and to identify configurations that may lead to congestion or delayed evacuation. It may also be useful for testing the impact of changes such as exit placement, corridor width, or occupancy levels. In this way, the approach provides a simple way to explore “what-if” scenarios without requiring detailed manual simulation setup.

9.9 Policy and Regulatory Implications

PeopleFlow is intended to complement, not replace, prescriptive compliance workflows. In practice, the framework can serve as an evidence layer for performance-based safety review by quantifying congestion hotspots, blocked-exit sensitivity, and clearance robustness under standardized stress scenarios. This supports transparent communication between designers, safety engineers, and authorities by adding reproducible scenario evidence to conventional rule-based checks, especially for high-occupancy facilities where geometry-induced bottlenecks can dominate risk. For jurisdictional review workflows, these outputs are suitable as supplementary technical evidence for Fire Marshal and Authority Having Jurisdiction (AHJ) discussions in performance-based evacuation assessments.

10 LIMITATIONS

This study has five main limitations. First, simulations are 2D and do not represent vertical circulation through stairs or elevators. A direct multi-floor extension would model each floor as a graph layer and encode stairs/elevators as inter-layer edges; the cost-function logic in Eq. (5) can then be evaluated on this augmented graph with additional vertical-transition penalties. Recent high-rise evacuation-elevator experiments further motivate this extension Mossberg et al. (2020). Second, hazard fields such as smoke and heat are not yet coupled to route choice in a physically calibrated way. Third, panic behavior is simplified, and human behavioral adaptation such as group cohesion and leader-following dynamics are approximated; group-level coordination and emergency communication effects are not explicitly modeled, which may influence evacuation dynamics in real deployments. Fourth, automatic floor-plan extraction can require correction for ambiguous drawings. Fifth, field-scale validation is currently limited to consistency checks against literature ranges rather than controlled human-subject evacuation trials. Supplementary airport/metro/office evaluations were expanded to full five-scenario coverage in this revision, but broader real-layout cohorts are still needed for stronger external generalization. These limitations define clear directions for future expansion and do not affect the comparative validity of the presented experiments.

11 FUTURE WORK

519 Future work will focus on (1) multi-floor 3D evacuation with staircase dynamics, (2) coupling
520 with fire and smoke propagation models, (3) reinforcement-learning-assisted adaptive routing
521 under hazards, and (4) live digital twin integration with sensor streams for online risk assessment.
522 For advanced routing-policy research, future extensions can also incorporate preference-aware
523 and compilation-based multi-agent pathfinding formulations to improve decision robustness in
524 constrained networks Ho and Nakadai (2022); Surynek (2022). We also plan to expand the open
525 benchmark set with additional plan categories such as hospital emergency floors, mall atria,
526 railway concourses, and auditorium layouts for stronger cross-domain generalization, and to
527 perform additional validation using controlled evacuation drills or publicly available trajectory
528 datasets.

529 **Practical Implications:** For near-term deployment, the most actionable recommendation is to
530 use PeopleFlow as a comparative screening tool during concept and retrofit stages: run structured
531 stress scenarios early, rank alternatives by congestion-risk indicators and clearance robustness,
532 and forward only the strongest candidates into detailed engineering review. This workflow reduces
533 redesign latency and improves traceability of safety decisions.

12 CONCLUSION

534 This paper presented a reproducible geometry-aware multi-agent framework for structured evalua-
535 tion of building evacuation safety performance. Across a 450-run core matrix and supplementary
536 parametric/generalization experiments, we showed substantial sensitivity of evacuation outcomes
537 to geometric topology and routing assumptions, including an observed 8.27x range in clearance
538 time and order-of-magnitude variation in peak density. Ablation, parameter-sensitivity, and
539 method-comparison experiments demonstrate that density-sensitive path-selection policies and
540 scenario design materially affect measured safety outcomes in structured safety-evaluation settings.

541 The work advances evacuation research by framing simulation as an auditable, reproducible
542 framework rather than a one-off prediction tool. The resulting computational workflow supports
543 transparent comparison for architects, safety engineers, and smart-city planners who need evidence-
544 based prioritization of safer design and policy options. Overall, PeopleFlow provides a practical
545 foundation for experiment-driven evacuation safety evaluation, supports data-driven infrastructure
546 design, and serves as a pre-deployment evidence layer for standards-aligned review, design iteration,
547 and operational preparedness assessment.

548 PeopleFlow supports early-stage evacuation planning by enabling rapid evaluation of archite-
549 ctural design alternatives prior to construction. The framework can assist safety engineers in
550 identifying congestion-prone regions, evaluating exit placement strategies, and testing emerge-
551 ncy routing policies within digital twin environments. Such simulation-driven design assessment
552 can contribute to safer high-occupancy infrastructure including transportation hubs, hospitals,
553 and commercial facilities. The workflow aligns with emerging digital twin methodologies for
554 safety-aware infrastructure modeling.

13 APPENDIX: DIGITAL TWIN INTEGRATION SCHEMA AND API HOOKS

555 To support deployment-oriented digital-twin pipelines, PeopleFlow session artifacts follow a stable
556 schema that separates run identity, geometry state, scenario parameters, and outcome metrics.

Table 22. PeopleFlow Artifact Schema for Digital-Twin Integration

Schema field	Description
run_id	Unique execution identifier with timestamp and seed tag
geometry_hash	Deterministic hash of processed floor-plan geometry
scenario_config	Occupancy, routing policy, blocked-exit flags, and behavioral settings
step_metrics	Time-series of local density, flow, queue length, and evacuation progress
aggregate_metrics	Clearance time, peak density indicator, bottleneck duration, exit utilization
provenance	Software version, parameter file checksum, and execution environment

557 Deployment-facing API hooks can expose this schema through standard endpoints:

- 558 • “/api/v2/simulations”
559 • “/api/v2/simulations/{id}/metrics”
560 • “/api/v2/simulations/{id}/artifacts”

561 In a live digital-twin loop, occupancy sensors update density estimates used in Eq. (5), while
562 archived artifacts provide audit-ready evidence for periodic safety review.

AUTHOR CONTRIBUTIONS

563 Likith Sai Kunchi: Conceptualization, Methodology, Software, Data Curation, Formal Analysis,
564 Visualization, Writing – original draft. [PhD Student Full Name]: Validation, Investigation,
565 Writing – review and editing. [Faculty Guide Full Name]: Supervision, Methodology Review,
566 Project Administration, Writing – review and editing.

CONFLICT OF INTEREST

567 The authors declare that the research was conducted in the absence of any commercial or financial
568 relationships that could be construed as a potential conflict of interest.

FUNDING

569 The author(s) received no specific funding for this work.

ACKNOWLEDGMENTS

570 The authors used GitHub Copilot and ChatGPT for coding assistance and language polishing
571 during manuscript preparation. All AI-assisted outputs were reviewed, technically verified, and
572 revised by the authors before inclusion. The authors take full responsibility for the content,
573 methodology, and technical accuracy of this work.

REFERENCES

- Alahi, A., Goel, K., Ramanathan, V., Robicquet, A., Fei-Fei, L., and Savarese, S. (2016). Social lstm: Human trajectory prediction in crowded spaces. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 961–971. doi:10.1109/cvpr.2016.110
- Batty, M. (2012). Smart cities, big data. *Environment and Planning B: Planning and Design* 39, 191–193
- Chen, K., Song, X., and Ren, X. (2021). Modeling social interaction and intention for pedestrian trajectory prediction. *Physica A: Statistical Mechanics and its Applications* 570, 125790. doi:10.1016/j.physa.2021.125790
- Chen, S. and Wang, J. (2018). Agent-based modeling of pedestrian evacuation considering panic and guidance. *Simulation Modelling Practice and Theory* 82, 10–23
- Chen, X. and Wang, J. (2021). Entropy-based crowd evacuation modeling with seeking behavior of social groups. *IEEE Access* 9, 4653–4664. doi:10.1109/access.2020.3048054
- Ding, Y., Zhang, Y., and Huang, X. (2023). Intelligent emergency digital twin system for monitoring building fire evacuation. *Journal of Building Engineering* 77, 107416. doi:10.1016/j.jobbe.2023.107416
- Duives, D. C., Daamen, W., and Hoogendoorn, S. P. (2013). State-of-the-art crowd motion simulation models. *Transportation Research Part C: Emerging Technologies* 37, 193–209
- Fruin, J. J. (1971). *Pedestrian Planning and Design* (New York, NY, USA: Metropolitan Association of Urban Designers and Environmental Planners)
- Gupta, A., Johnson, J., Fei-Fei, L., Savarese, S., and Alahi, A. (2018). Social gan: Socially acceptable trajectories with generative adversarial networks. In *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2255–2264. doi:10.1109/cvpr.2018.00240
- Gwynne, S., Galea, E., Owen, M., Lawrence, P., and Filippidis, L. (1999). A review of the methodologies used in the computer simulation of evacuation from the built environment. *Building and Environment* 34, 741–749
- Haghani, M. (2020). Empirical methods in pedestrian, crowd and panic dynamics: A review. *Safety Science* 129, 104841
- Helbing, D., Farkas, I., and Vicsek, T. (2000). Simulating dynamical features of escape panic. *Nature* 407, 487–490
- Helbing, D. and Molnar, P. (1995). Social force model for pedestrian dynamics. *Physical Review E* 51, 4282–4286
- Ho, F. and Nakadai, S. (2022). Preference-based multi-objective multi-agent path finding. *Autonomous Agents and Multi-Agent Systems* 37. doi:10.1007/s10458-022-09593-3
- Hurley, M. J., Gottuk, D., Hall, J. R. J., Harada, K., Kuligowski, E. D., Puchovsky, M., et al. (eds.) (2016). *SFPE Handbook of Fire Protection Engineering* (New York, NY, USA: Springer), 5 edn.
- Jahangir, M. F., Kamari, A., and Schultz, C. P. L. (2025). Unified and interoperable digital twin models for emergency evacuation using building simulation identity cards. *Automation in Construction* 175, 106225. doi:10.1016/j.autcon.2025.106225
- Jiang, L., Shi, J., Wang, C., and Pan, Z. (2023). Intelligent control of building fire protection system using digital twins and semantic web technologies. *Automation in Construction* 147, 104728. doi:10.1016/j.autcon.2022.104728

- Kirchner, A. and Schadschneider, A. (2002). Simulation of evacuation processes using a bionics-inspired cellular automaton model for pedestrian dynamics. *Physica A: Statistical Mechanics and its Applications* 312, 260–276.
- Kuligowski, E. D. (2013). Predicting human behavior during fires. *Fire Technology* 49, 101–120.
- Li, X., Liu, Y., Wang, K., and Wang, F.-Y. (2020). A recurrent attention and interaction model for pedestrian trajectory prediction. *IEEE/CAA Journal of Automatica Sinica* 7, 1361–1370. doi:10.1109/jas.2020.1003300
- Lopez-Carmona, M. A. and Paricio Garcia, A. (2022). Adaptive cell-based evacuation systems for leader-follower crowd evacuation. *Transportation Research Part C: Emerging Technologies* 140, 103699. doi:10.1016/j.trc.2022.103699
- Lovreglio, R., Gonzalez, V., Feng, Z., Amor, R., Spearpoint, M., Thomas, J., et al. (2020). Prototyping virtual reality serious games for building earthquake preparedness: The auckland city hospital case study. *Advanced Engineering Informatics* 32, 100900.
- Mohamed, A., Qian, K., Elhoseiny, M., and Claudel, C. (2020). Social-stgcn: A social spatio-temporal graph convolutional neural network for human trajectory prediction. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 14412–14420. doi:10.1109/cvpr42600.2020.01443
- Mossberg, A., Nilsson, D., and Andree, K. (2020). Unannounced evacuation experiment in a high-rise hotel building with evacuation elevators: A study of evacuation behaviour using eye-tracking. *Fire Technology* 57, 1259–1281. doi:10.1007/s10694-020-01046-1
- Nishinari, K., Kirchner, A., Namazi, A., and Schadschneider, A. (2004). Extended floor field CA model for evacuation dynamics. *IEICE Transactions on Information and Systems* E87-D, 726–732.
- Osama, Z. (2024). The digital twin framework: A roadmap to the development of user-centred digital twin in the built environment. *Journal of Building Engineering* 98, 111081. doi:10.1016/j.jobe.2024.111081
- Ozturk, G. B. (2021). Digital twin research in the AECO-FM industry. *Journal of Building Engineering* 40, 102730. doi:10.1016/j.jobe.2021.102730
- Pan, X., Han, C. S., Dauber, K., and Law, K. H. (2007). A multi-agent based framework for the simulation of human and social behaviors during emergency evacuations. *AI & Society* 22, 113–132.
- Pellegrini, S., Ess, A., Schindler, K., and van Gool, L. (2009). You'll never walk alone: Modeling social behavior for multi-target tracking. In *2009 IEEE 12th International Conference on Computer Vision*. 261–268. doi:10.1109/iccv.2009.5459260
- Ronchi, E. and Gwynne, S. (2015). Data collection for evacuation model validation. *Fire Safety Journal* 72, 137–151.
- Ronchi, E. and Nilsson, D. (2013). Fire evacuation in high-rise buildings: a review of human behaviour and modelling research. *Fire Science Reviews* 2, 1–21.
- Saleh, K. (2022). Pedestrian trajectory prediction for real-time autonomous systems via context-augmented transformer networks. *Sensors* 22, 7495. doi:10.3390/s22197495
- Salzmann, T., Ivanovic, B., Chakravarty, P., and Pavone, M. (2020). Trajectron++: Dynamically-feasible trajectory forecasting with heterogeneous data. In *Computer Vision – ECCV 2020* (Springer). 683–700. doi:10.1007/978-3-030-58523-5_40
- Still, G. K. (2014). *Introduction to Crowd Science* (Boca Raton, FL, USA: CRC Press)

- 660 Sun, Q. and Turkan, Y. (2020). A BIM-based simulation framework for fire safety management
661 and investigation of the critical factors affecting human evacuation performance. *Advanced*
662 *Engineering Informatics* 44, 101093. doi:10.1016/j.aei.2020.101093
- 663 Surynek, P. (2022). Problem compilation for multi-agent path finding: A survey. In *Proceedings*
664 *of the Thirty-First International Joint Conference on Artificial Intelligence*. 5615–5622. doi:10.
665 24963/ijcai.2022/783
- 666 Wang, R., Song, X., Hu, Z., and Cui, Y. (2023). Spatio-temporal interaction aware and trajectory
667 distribution aware graph convolution network for pedestrian multimodal trajectory prediction.
668 *IEEE Transactions on Instrumentation and Measurement* 72, 1–11. doi:10.1109/tim.2022.
669 3229733
- 670 Wang, X., Chraibi, M., Chen, J., Li, R., and Ma, J. (2022). Modeling boundedly rational route
671 choice in crowd evacuation processes. *Safety Science* 147, 105590. doi:10.1016/j.ssci.2021.105590
- 672 Wehbe, R. and Shahrou, I. (2021). A BIM-based smart system for fire evacuation. *Future*
673 *Internet* 13, 221. doi:10.3390/fi13090221
- 674 Weidmann, U. (1993). *Transporttechnik der Fussgänger: Transporttechnische Eigenschaften*
675 *des Fussgängerverkehrs* (Zurich, Switzerland: ETH Zurich, Institut für Verkehrsplanung,
676 Transporttechnik, Strassen- und Eisenbahnbau)
- 677 Yano, R. (2018). Effect of form of obstacle on speed of crowd evacuation. *Physical Review E* 97.
678 doi:10.1103/physreve.97.032319
- 679 Zhang, C., Lu, Y., and Tao, F. (2021). Digital twins for smart cities: A state-of-the-art review.
680 *Engineering* 7, 1715–1730
- 681 Zheng, X., Zhong, T., and Liu, M. (2009). Modeling crowd evacuation of a building based on
682 seven methodological approaches. *Building and Environment* 44, 437–445
- 683 Zhou, H., Ren, D., Xia, H., Fan, M., Yang, X., and Huang, H. (2021). Ast-gnn: An attention-based
684 spatio-temporal graph neural network for interaction-aware pedestrian trajectory prediction.
685 *Neurocomputing* 445, 298–308. doi:10.1016/j.neucom.2021.03.024