

PeopleFlow: A Reproducible Geometry-Aware Multi-Agent Framework for Comparative Building Evacuation Safety Evaluation

Likith Sai Kunchi^{1,*}

¹*B.Tech CSE (SCOPE), VIT (Vellore Institute of Technology), India*¹*B.Tech CSE (SCOPE), VIT (Vellore Institute of Technology), India*

Correspondence*:

Likith Sai Kunchi

likithsai.kunchi@gmail.comLikith Sai Kunchi

likithsai.kunchi@gmail.com

2 ABSTRACT

3 Evaluating evacuation performance in complex buildings requires analysis methods that capture
4 interactions between spatial configuration, pedestrian behavior, and operational disruptions.
5 This paper introduces PeopleFlow, a layout-informed computational framework for structured
6 evacuation-safety comparison across diverse architectural layouts. The workflow converts floor-
7 plan representations into navigable simulation environments and applies continuous-space
8 multi-agent modeling based on force-driven pedestrian dynamics.

9 PeopleFlow reports interpretable indicators including total clearance duration, localized
10 congestion intensity, bottleneck persistence, and exit-utilization balance. To enable consistent
11 experimentation, the framework uses parameterized scenario configurations and reproducibility-
12 oriented execution records. The evaluation covers a structured matrix of nine layout types across
13 five operational conditions (450 core runs), plus 210 controlled supplementary/sensitivity runs for
14 real-layout generalization and sensitivity to exit width, desired speed, and occupancy scaling.

15 Results show strong dependence of evacuation outcomes on geometric topology and routing
16 strategy, with clearance-time differences exceeding an eightfold range and substantial congestion
17 variation across layouts. The primary contribution is a reproducible engineering framework
18 for relative safety prioritization across heterogeneous architectural configurations, rather than
19 deterministic prediction of individual trajectories. The framework supports what-if analysis during
20 early design and provides a computational basis for integration with digital twin workflows for
21 high-occupancy safety planning.

22 **Keywords:** Crowd simulation, evacuation analytics, evacuation modelling, digital twin, multi-agent modelling, multi-agent simulation,
23 smart infrastructure, building evacuation safety, reproducible pipeline.

24 sectionIntroduction

25 Building evacuation performance evaluation is a high-impact engineering task where static
26 compliance checks often miss dynamic crowd effects under disruption. Historical emergencies
27 consistently show that bottleneck formation, uneven exit utilization, and localized panic response
28 can dominate outcomes even in code-compliant facilities. Conventional hand calculations remain

useful for baseline checks, but they provide limited support for comparing design alternatives under multi-hazard and operational uncertainty Hurley et al. (2016); Ronchi and Nilsson (2013).

Recent high-density infrastructure such as airports, metro systems, hospitals, and large commercial complexes introduces complex evacuation dynamics that cannot be fully evaluated using static compliance rules alone.

The need for early-stage design validation has increased with higher occupancy density in modern infrastructure. Although simulation enables repeatable structured comparison before deployment, the central challenge is integrating architectural input, scenario control, execution, and interpretable outputs within one coherent computational workflow.

The present research introduces PeopleFlow as a layout-informed evacuation efficiency framework with traceable computational pipeline execution. Rather than proposing a new pedestrian-physics law in isolation, the study develops a unified computational pipeline for robust scenario-based evaluation across heterogeneous layouts.

The main contributions are summarized as follows:

- A reproducible floor-plan-aware simulation workflow linking layout ingestion, scenario execution, and artifact-backed reporting.
- A structured cross-layout evacuation comparison across 660 simulations (450 core + 210 supplementary/sensitivity) spanning diverse building archetypes.
- A density-aware routing evaluation framework under structured stress scenarios, supported by ablation and sensitivity analysis.

While prior evacuation simulators emphasize pedestrian behavioral realism or scenario-specific visualization, they largely lack automated floor-plan ingestion, reproducible batch execution, and artifact-native cross-layout benchmarking within a single coherent pipeline. PeopleFlow addresses this gap directly. To the best of our knowledge, this is the first openly reproducible framework that couples automated floor-plan ingestion, stress-scenario batch execution, and artifact-native auditing for cross-layout evacuation-safety ranking – without requiring manual geometry preparation or proprietary tooling. The contribution is methodological and infrastructural: a transparent comparative-evaluation pipeline that enables architects, safety engineers, and digital-twin integrators to systematically prioritize safer design alternatives under structured stress conditions.

1 RELATED WORK

1.1 Crowd Modeling Approaches

Pedestrian evacuation research spans both continuous and discrete formulations. In force-based formulations, evacuees can be represented as self-driven particles whose motion evolves through interaction potentials capturing repulsion, attraction, and directional preference; this family remains a physically interpretable baseline for dense-crowd analysis Helbing and Molnar (1995); Helbing et al. (2000). Cellular automata and floor-field variants offer efficient large-scale approximations but may introduce directional discretization artifacts near bottlenecks and tight corners Nishinari et al. (2004); Kirchner and Schadschneider (2002). Comparative studies

67 indicate that model suitability is task-dependent; for complex indoor geometries, continuous-
68 space approaches often preserve turning behavior and local interaction structure more naturally
69 Zheng et al. (2009); Duives et al. (2013); Still (2014). Recent pedestrian-trajectory literature
70 further improves interaction modeling using benchmark datasets and social-context forecasting
71 architectures, including ETH/UCY-aligned methods based on recurrent, graph, and generative
72 formulations Pellegrini et al. (2009); Alahi et al. (2016); Gupta et al. (2018); Salzmann et al.
73 (2020); Mohamed et al. (2020); Zhou et al. (2021); Chen et al. (2021); Li et al. (2020); Wang et al.
74 (2023); Saleh (2022).

75 1.2 Evacuation Simulation Systems

76 Commercial and academic evacuation platforms generally cover geometry authoring, route
77 simulation, and report generation, but they differ substantially in automation depth, extensibility,
78 and experiment traceability Gwynne et al. (1999); Ronchi and Gwynne (2015); Lovreglio et al.
79 (2020). Behavior-focused studies further show that guidance logic, uncertainty handling, and
80 panic-modulated decisions can materially alter evacuation outcomes Pan et al. (2007); Chen and
81 Wang (2018). In many practical settings, heavy manual preprocessing and proprietary tooling still
82 constrain transparent cross-layout benchmarking.

83 1.3 Digital Twin Safety Analysis

84 Digital twin studies in the built environment increasingly target safety and emergency operations.
85 Prior work supports simulation-coupled twins for risk-aware planning, yet complete pipelines
86 from floor-plan ingestion to repeatable quantitative evacuation indicators remain limited in many
87 deployments Batty (2012); Zhang et al. (2021); Ozturk (2021); Osama (2024). Recent BIM-tied
88 fire-safety and emergency-digital-twin systems also demonstrate practical monitoring, evacuation
89 analytics, and control integration, reinforcing the need for end-to-end reproducible computational
90 workflows Sun and Turkan (2020); Wehbe and Shahrour (2021); Ding et al. (2023); Jiang et al.
91 (2023); Jahangir et al. (2025). These trends remain active in 2024–2026 deployment-oriented safety
92 literature, with emphasis shifting from isolated twin prototypes toward interoperable, reproducible
93 engineering workflows that can be independently rerun and audited.

94 1.4 Research Gap and Positioning

95 The gap addressed in this paper is methodological: reproducible, layout-informed evacuation
96 efficiency comparison across diverse layouts and stress scenarios. PeopleFlow is positioned as a
97 reproducible framework rather than a claim of absolute predictive certainty. Its novelty lies in
98 coupling geometry extraction, session-managed simulation, scenario ablation, and interpretable
99 reporting within one coherent computational pipeline. Unlike many existing evacuation simulators
100 that require manual geometry preparation or proprietary workflows, PeopleFlow emphasizes
101 reproducibility through artifact-backed records, enabling transparent cross-layout comparison
102 under consistent scenario definitions. Recent 2024–2026 deployment-oriented safety literature
103 has accelerated the shift from isolated simulation prototypes toward interoperable, reproducible
104 engineering workflows. Jahangir et al. (2025) demonstrated unified digital twin
105 models for emergency evacuation using building simulation identity cards, while Ding et al. (2023)
106 introduced intelligent emergency digital twin systems for real-time fire evacuation
107 monitoring. Despite these advances, no existing open framework couples automated floor-plan

Table 1. Comparison with Existing Evacuation Tools

Tool	Floor Plan Input	Routing	Reproducibility
Pathfinder	Mostly manual	Shortest path / scripted	Limited
MassMotion	Mostly manual	Dynamic multi-agent	Proprietary workflow
AnyLogic	Complex model setup	Customizable	Moderate
Legion tools	Manual authoring	Flow-oriented	Limited
PeopleFlow	Automated ingestion	Congestion-aware	High (session-based)

108 ingestion, stress-scenario batch execution, and artifact-native auditing for reproducible cross-
 109 layout safety ranking – the precise gap PeopleFlow addresses. Unlike commercial tools such
 110 as Pathfinder and MassMotion which require manual geometry preparation and offer limited
 111 reproducibility, PeopleFlow provides session-managed, auditable experiment artifacts enabling
 112 transparent cross-layout comparison under consistent scenario definitions. A concise capability
 113 comparison with representative tools is provided in Table 1.

2 METHODOLOGY

114 2.1 Workflow Pipeline

115 PeopleFlow follows a deterministic workflow with four stages: (1) floor-plan ingestion, (2)
 116 geometry extraction and navigation graph generation, (3) scenario-based simulation execution,
 117 and (4) artifact-backed analytics. Each run is stored with configuration metadata, seeds, and
 118 per-step metrics so results are replayable and auditable. The step-level simulation procedure is
 119 summarized in Fig. 4.

120 2.2 Reproducibility Design

121 Each experiment is represented by a session artifact that captures:

- 122 • Input layout identifier and geometry hash.
- 123 • Scenario parameters (occupancy, blocked exits, routing policy, panic level).
- 124 • Random seed and simulation step size.
- 125 • Aggregated outputs and raw trajectory summaries.

126 This structure allows direct reruns and supports batch-level statistical comparisons. The
 127 modular organization also enables independent evaluation of routing logic without modifying
 128 geometric preprocessing components. Simulation configurations and synthetic benchmark
 129 layouts are packaged as reproducible experiment artifacts and can be released to support
 130 future comparative studies. The implementation and reproducibility assets are maintained at
 131 <https://github.com/1206likith/PeopleFlow>. Experiment artifacts allow deterministic reruns and




Figure 1. Agent movement logic used in each simulation step.

structured comparison across layout categories. A 30-second multimedia supplement demonstrating the end-to-end ingestion-to-analytics workflow is included in the submission artifacts.

2.3 Agent Decision Logic

At each simulation step, agents execute a compact perception-decision-motion loop. This structure ensures behavioral consistency while allowing routing-policy variation to be isolated during controlled ablation experiments. Local density and queue conditions are sensed first, then exit costs are updated using Eq. (5). The lowest-cost exit is selected, social-force dynamics are integrated, and state variables are updated for the next step.

2.4 Computational Complexity

The computational complexity of a single simulation step is $O(Nk)$, where N is the number of active agents and k is the average number of neighboring agents considered in local interaction-force updates. Under practical density regimes, k remains locally bounded, giving near-linear scaling with population size for each step. For batch execution, total workload scales as $O(RNS)$, where R is the number of runs and S is the number of simulation steps per run. For a layered 3D extension with inter-floor connectors, the step cost becomes $O(Nk + M)$ where M is the number of active inter-layer connector evaluations, preserving near-linear practical scaling for sparse connector graphs. This complexity profile motivates the reproducible batch workflow used in our large-matrix experiments. Simulation complexity therefore scales approximately linearly with agent population size under bounded interaction-radius assumptions. Measured wall-clock trends (Table 5) are consistent with this analysis and show predictable runtime growth with scenario and layout complexity under fixed hardware conditions.




Figure 2. Social-force interaction components used for each movement update.

Table 2. Scenario Definitions

Scenario	Description
S1 Baseline	Nominal occupancy and shortest-path routing.
S2 HighOcc	Increased occupancy stress to test crowding sensitivity.
S3 Blocked	One major exit disabled to evaluate rerouting resilience.
S4 Routing	Congestion-aware routing policy activated.
S5 Panic	Stochastic panic variation in desired speed and interactions.

153 2.5 Layered 3D Transition Logic

154 For a layered multi-floor extension, we define an inter-layer transition cost for moving from floor
 155 level z_1 to z_2 through connector mode m (stairs or elevator):

$$C_{\text{layer}}(z_1, z_2, m, t) = \eta_d |z_2 - z_1| + \eta_t \tau_m + \eta_q q_m(t), \quad (1)$$

156 where τ_m is connector traversal delay and $q_m(t)$ is a time-varying queue penalty at connector m .
 157 This term is added to Eq. (5) when inter-floor routing alternatives are available.

158 2.6 Scenario Matrix

159 The study uses a five-scenario matrix reflecting realistic stress transitions from nominal operation
 160 to disruption. Table 2 summarizes the definitions.




Figure 3. Layer-aware routing extension.

3 PEOPLEFLOW ARCHITECTURE

161 3.1 System Components

162 The architecture consists of geometry ingestion, simulation kernel, metrics engine, and reporting
 163 layer. Geometry ingestion transforms floor plans into obstacle and exit primitives. The simulation
 164 kernel executes force-based motion and route updates. The metrics engine computes clearance,
 165 density, bottleneck, and utilization indicators. The reporting layer produces figures and tables
 166 for comparative studies. Figure 5 and Figure 6 summarize these system-level and workflow-level
 167 connections.

168 The architecture emphasizes separation between geometry processing, simulation execution, and
 169 artifact-based reporting.

170 3.2 Session and Artifact Model

171 A session-centric design is used to support scientific traceability. For each run, PeopleFlow
 172 stores configuration, random seed, outcome metrics, and derived plots. This explicit artifact chain
 173 enables transparent comparisons between methods and scenario variants.




Figure 4. PeopleFlow simulation loop.

174 3.3 Analytics Outputs

175 Standard outputs include: clearance-time distribution, peak-density maps, bottleneck duration
 176 traces, exit-load statistics, and method-comparison summaries. These outputs are generated from
 177 reproducible logs instead of ad hoc post-processing.

4 MATHEMATICAL MODEL

178 4.1 Social Force Dynamics

179 Each agent i is modeled by a second-order force balance:

$$m_i \frac{d\mathbf{v}_i}{dt} = \mathbf{f}_i^{\text{drv}} + \sum_{j \neq i} \mathbf{f}_{ij}^{\text{int}} + \sum_W \mathbf{f}_{iW}^{\text{wall}}. \quad (2)$$

180 The social force formulation was selected due to its interpretability, continuous-space representation,
 181 and extensive validation in pedestrian-dynamics literature, making it suitable for comparative




Figure 5. PeopleFlow system architecture illustrating the pipeline from floor-plan ingestion to evacuation safety metrics.




Figure 6. Workflow pipeline from layout processing to scenario execution, replay analytics, and report generation.

182 evacuation analysis. The driving force is:

$$\mathbf{f}_i^{\text{drv}} = m_i \frac{v_i^0 \mathbf{e}_i - \mathbf{v}_i}{\tau_i}, \quad (3)$$

183 where v_i^0 is desired speed, $\mathbf{e}_i$ is desired direction, and τ_i is relaxation time. Here, τ_i acts as a
 184 reaction-time parameter controlling how quickly agents adapt their actual velocity toward desired
 185 motion. In the implemented discrete kernel, force integration uses normalized agent mass $m_i = 1$,
 186 update step $\Delta t = 0.2$ s, and social-force gain 0.2, which together define the effective adaptation
 187 timescale.

Table 3. Routing Cost Parameters Used for Reproducibility

Parameter	Value used in experiments
α (distance weight)	1.0 (baseline); sensitivity checks in the 0.8–1.2 range
β (density weight)	2.0 (baseline); sensitivity checks in the 1.5–2.5 range
γ (hazard weight)	0.0 for core/supplementary runs (hazard fields disabled)
Queue perception radius	5.0 m around each exit
Stochastic softmax temperature T	15.0

4.2 Density-Flow Relation

At a macroscopic level, local flow follows:

$$q = \rho v, \quad (4)$$

where q denotes pedestrian flow rate, ρ denotes local pedestrian concentration, and v represents average movement velocity. This relation is used for consistency checks between simulation outputs and empirical benchmarks Fruin (1971); Weidmann (1993).

4.3 Exit Selection Cost Function

For congestion-aware routing, the exit cost at time t is:

$$C_e(t) = \alpha d_e(t) + \beta \rho_e(t) + \gamma h_e(t), \quad (5)$$

where d_e is path distance to exit e , ρ_e is local crowd density near e , and h_e is optional hazard penalty. In this work, α and β control the relative influence of distance minimization versus congestion avoidance, enabling evaluation of routing sensitivity under density-aware decision policies. The routing formulation balances geometric shortest-path preference with local density avoidance, enabling dynamic redistribution of agents across available exits under congestion conditions. In the core and supplementary experiments, congestion-aware routing uses $\alpha = 1.0$, $\beta = 2.0$, and $\gamma = 0.0$ (hazard term disabled), with a local queue-perception radius of 5.0 m; stochastic routing uses the same base cost with softmax temperature $T = 15.0$.

4.4 Statistical Metrics

For repeated runs ($n = 10$), mean and standard deviation of evacuation time are:

$$\bar{T} = \frac{1}{n} \sum_{k=1}^n T_k, \quad (6)$$

$$\sigma_T = \sqrt{\frac{1}{n} \sum_{k=1}^n (T_k - \bar{T})^2}. \quad (7)$$

Table 4. Simulation Parameters

Parameter	Value
Time step (Δt)	0.2 s
Nominal desired speed	1.2 m/s
Agent radius	0.3 m
Social-force mass normalization (m_i)	1.0 (dimensionless, normalized integration)
Relaxation parameter (τ_i)	Implicit in discrete update ($\Delta t = 0.2$ s, force gain 0.2)
Runs per configuration	10
Routing policies	nearest, congestion-aware, guided/stochastic variants
Scenarios	S1 to S5 (Table 2)
Total runs (core + supplements)	660
Execution setup	Local Python backend workflow on Windows workstation

206 Exit load imbalance is measured as:

$$I_{\text{exit}} = \frac{1}{2N} \sum_{e=1}^E \left| n_e - \frac{N}{E} \right|, \tag{8}$$

207 where n_e is evacuated count via exit e , N is total evacuated agents, and E is number of usable
208 exits.

5 EXPERIMENTAL SETUP

209 5.1 Simulation Parameters

210 The core study executes 450 runs (9 layouts $\times$ 5 scenarios $\times$ 10 repetitions). To strengthen
211 generalization and engineering interpretability, we additionally conduct 210 controlled runs:
212 150 supplementary real-layout runs (airport terminal, metro station, office building; full S1–S5
213 coverage), 30 exit-width comparison runs, and 30 desired-speed sensitivity runs. Key simulation
214 parameters are shown in Table 4. Nine layouts were selected to balance geometric diversity with
215 computational feasibility while enabling statistically meaningful cross-layout comparison under a
216 fixed 10-seed protocol and reproducible runtime budget. Parameter ranges were selected based
217 on commonly reported pedestrian-dynamics studies describing typical free-flow speeds between
218 1.2–1.5 m/s and density thresholds consistent with established fundamental-diagram behavior.

219 5.2 Implementation Details

220 The simulation engine is implemented in Python with NumPy-based state updates and
221 deterministic seed control for reproducibility. Geometry preprocessing uses vector-edge extraction
222 and obstacle/exit parsing before navigation-graph construction. Experiments were executed on a
223 workstation-class environment (Intel i7 CPU class, 16 GB RAM), with artifact logging enabled
224 for run replay and statistical post-processing. The floor-plan ingestion pipeline follows an explicit
225 computer-vision sequence: adaptive and Otsu thresholding for linework isolation, Canny edge

Table 5. Measured Wall-Clock Runtime and Scalability Evidence (5 Runs per Layout)

Layout	Wall-clock (s)	Mean steps	Completion rate
Airport terminal	10.19 ± 1.48	179.2	0.60
Metro station	1.48 ± 0.11	154.8	1.00
Office building	0.95 ± 0.03	147.8	1.00
Plan3 (core set)	1.07 ± 0.18	1254.8	1.00
Layered 3D extension [†]	$1.15 \times - 1.35 \times$ 2D runtime	—	—

extraction (50/150 thresholds), and probabilistic Hough line detection ($\rho = 1$, $\theta = \pi/180$) with scale-adaptive thresholds. Exit candidates are then filtered by width constraints (minimum 10 px and maximum 12% of the shortest image dimension) and proximity checks to wall or boundary segments before deduplication.

5.3 Validity Controls and Reporting Protocol

Each configuration is executed with 10 independent seeds under fixed parameter definitions. We report mean $\pm$ standard deviation in all compact comparison tables and use identical logging pipelines across core and supplementary cohorts. For selected engineering sweeps, confidence intervals can be computed as $\bar{T} \pm 1.96 \sigma_T / \sqrt{n}$ with $n = 10$. Supplementary airport/metro/office evaluations use full five-scenario coverage under the same protocol to test cross-geometry transfer under aligned stress definitions. All experiments were executed under identical numerical integration settings to ensure consistency across scenario comparisons.

5.4 Runtime and Scalability Evidence

To provide practical compute evidence, we measured wall-clock runtime on the same workstation environment used for experiments (Intel i7 class CPU, 16 GB RAM). Each layout was evaluated for 5 repeated baseline runs (80 agents, $\Delta t = 0.2$ s). For Plan3, the baseline scalability rerun was executed at $T_{\max} = 350$ s using S1 shortest-path routing to avoid timeout-only artifacts. Table 5 reports observed runtime statistics.

[†]Projection derived from the $O(Nk + M)$ extension analysis with sparse inter-floor connector checks; this row is a complexity-based estimate, not a measured runtime. Figure 19 visualizes the same runtime evidence with uncertainty bars and completion-rate trends for direct comparative interpretation.

5.5 Workflow Setup Effort Comparison

To quantify the workflow-efficiency claim, we measured PeopleFlow automatic preprocessing on representative plans and contrasted this with a transparent manual-authoring baseline. Because direct proprietary-tool benchmarking is not executable in this environment, manual setup effort is reported as a practitioner range from workflow-oriented evacuation-tool literature and engineering practice Gwynne et al. (1999); Ronchi and Gwynne (2015).

5.6 Floor-Plan Processing and Simulation Environment

Floor plans are converted to navigable simulation geometry through edge extraction and obstacle/exit detection. Figure 7 shows the conversion workflow, while Figure 8 and Figure 9 show representative processed and runtime views.

Table 6. Setup-Effort Comparison for Simulation Preparation

Workflow	Human steps	Operator time	Automated preprocessing
PeopleFlow	Upload + optional exit verification	1–3 min	0.51–67.58 s (mean 27.01 s)
Manual authoring baseline	Geometry tracing + exit placement + scenario setup	30–120 min (operator dependent)	Not automated




Figure 7. Floor-plan to simulation-geometry conversion workflow.

257 All simulation seeds, configuration parameters, and layout identifiers are stored in session
258 artifacts to enable deterministic reruns and statistical verification.

6 CASE STUDY FLOOR PLANS

259 We evaluate a core set of nine layouts available in the project floor-plan set: “Academic_Plan,”
260 “Hospital_Plan,” “Hotel_Plan,” “Library_Plan,” “Mall_Plan,” “Supermarket_Plan,” and three
261 additional complex plans (“Plan1,” “Plan2,” “Plan3”). These provide heterogeneous layouts
262 including academic, hospital, mall, and corridor-dominant structures. To improve real-
263 world diversity, we also include a supplementary cohort with “Airport_Terminal_Plan,”
264 “Metro_Station_Plan,” and “Office_Building_Plan” under full five-scenario coverage for cross-layout
265 generalization checks. The analyzed layout set and its geometric characteristics are summarized in

Table 7. Case Study Layouts

Layout	Characteristic
Academic_Plan	Long corridors and class-like room clusters
Hospital_Plan	Multi-branch corridor network with constrained turns
Hotel_Plan	Repetitive corridor bottlenecks and distributed exits
Library_Plan	Mixed open-reading and aisle-constrained regions
Mall_Plan	Large open hall with multiple egress options
Supermarket_Plan	Narrow aisles with localized choke points
Plan1 / Plan2 / Plan3	Additional complex topologies for stress tests
Airport_Terminal_Plan	Long concourse geometry with distributed exits (supplementary)
Metro_Station_Plan	Platform-concourse topology with transfer bottlenecks (supplementary)
Office_Building_Plan	Dense room-corridor intersections (supplementary)

Table 7. These layout categories represent commonly studied evacuation archetypes in fire safety and crowd engineering literature. The portfolio covers representative university-like blocks, hospital emergency-floor topology, shopping-mall atrium behavior, metro-station concourse circulation, and airport-terminal flow patterns for practical applicability.

7 RESULTS AND STATISTICAL ANALYSIS

7.1 Main Multi-Layout Results

Table 8 reports means and standard deviations for evacuation time and peak density across all 45 layout-scenario pairs. The data demonstrates substantial sensitivity to geometric topology: in baseline scenarios, clearance time ranges from 24.18 s (“Mall_Plan”) to 182.68 s (“Academic_Plan”), while blocked-exit stress drives extreme values up to 350.0 s (“Plan3”).

Low variance across repeated simulation seeds indicates stability of comparative layout rankings under stochastic variability. Beyond absolute values, the key statistical pattern is multi-seed ranking stability: layouts with corridor concentration remain systematically slower and denser than open layouts under S1/S2/S3. Across the 45 core layout-scenario cells (10 seeds each), 95% confidence intervals preserve the same high-risk and low-risk layout groups, indicating that comparative conclusions are robust to stochastic variation even when absolute outcomes shift by seed. This supports the intended use of the framework for relative safety prioritization rather

than exact second-level trajectory prediction. The large variation across layouts highlights the dominant influence of spatial topology compared to behavioral parameter variation alone.

Plan3 S3 Blocked previously reported a hard-timeout result of 200.0s due to an insufficiently set simulation ceiling. Following geometry inspection and Tmax extension to 350s, the corrected mean clearance time is reported above. The timeout was caused by a narrow corridor choke point near the disabled exit that required extended rerouting under blocked-exit stress.

7.2 Ablation Study

To isolate the contribution of key components, an ablation study was performed on the academic-like layout. Results in Table 9 confirm that congestion-aware routing is a dominant factor: removing it increases mean clearance from 7.24 s to 39.42 s in the controlled benchmark and raises mean density from 3.96 to 28.58. This trend is consistent with prior route-choice and adaptive-leader evacuation analyses in complex building evacuations Wang et al. (2022); Lopez-Carmona and Paricio Garcia (2022).

7.3 Method-Level Aggregate Comparison

A complementary method benchmark over corridor-centered stress tests is summarized in Table 10. Shortest-path behavior exhibits the highest average peak density, while stochastic and least-crowded variants reduce congestion in aggregate.

7.4 Comparison with Publicly Used Baseline Assumptions

To make method-level value more explicit, Table 11 summarizes behavioral differences among publicly used baseline assumptions in evacuation studies and practice.

7.5 Blocked-Exit Stress Comparison Against Baselines

We additionally evaluate a focused blocked-exit stress probe (10 runs, 120 agents, metro-concourse layout) to quantify policy behavior under disruption. Results are shown in Table 12.

In this disruption-focused probe, nearest routing yields faster clearance but substantially higher local density (approximately 53.7% higher peak-density indicator), while congestion-aware routing sacrifices clearance time to reduce hotspot intensity and improve crowd-distribution safety margins.

7.6 Supplementary Real-Layout Generalization

To strengthen cross-layout generalization, we expanded the supplementary real-layout cohort from baseline-only evaluation to full five-scenario coverage following the same protocol applied to the core matrix. Table 13 (expanded) reports results across S1–S5 for airport terminal, metro station, and office building layouts. Results confirm that scenario sensitivity patterns observed in the core matrix – particularly the dominant effect of blocked-exit stress and the congestion-reduction benefit of density-aware routing – generalize consistently across these real-world archetypes. The office building layout shows the strongest sensitivity to blocked exits due to dense room-corridor convergence, while the airport terminal exhibits the highest variance under panic conditions owing to its long concourse geometry.

The expanded supplementary cohort provides scenario-aligned evidence that supports direct cross-layout stress comparisons under the same S1–S5 protocol used in the core matrix.

7.7 Engineering Comparison: Effect of Exit Width

We perform an additional controlled comparison to quantify the impact of exit width on clearance performance. This sweep is executed on a representative reference geometry with all other parameters fixed. Table 14 shows that increasing width reduces evacuation time nonlinearly, with substantial gains between 1.0 m and 1.5 m.

7.8 Sensitivity Analysis: Desired Speed

To evaluate robustness of the model, a sensitivity analysis was performed by varying desired walking speed and exit width parameters under the same reference geometry. Table 15 shows that higher desired speeds improve clearance time but can increase local density peaks near constrained exits. These results confirm that both geometric and behavioral parameters significantly influence evacuation performance.

To isolate routing-cost sensitivity, Fig. 11 reports a grid sweep over α and β using recorded experiment artifacts. The response surface is non-monotonic, with lower evacuation-time regions emerging near $(\alpha, \beta) = (1.2, 2.0)$ and $(1.4, 4.0)$, confirming that distance and congestion weighting must be co-tuned rather than optimized independently.

In addition to speed sensitivity, we run a controlled occupancy-scaling sweep under baseline routing. Table 16 shows nonlinear growth in evacuation time and peak density as occupancy increases, providing a statistically grounded stress test for deployment planning.

7.9 Global Sensitivity Analysis (Sobol-Style)

To complement one-factor sweeps, we apply a Sobol-style variance ranking over the controlled factors (occupancy, exit width, and desired speed) using the same artifact-backed pipeline. The proxy total-order ranking indicates that occupancy and exit width dominate evacuation-time variance in the tested ranges, while desired speed contributes a smaller but non-negligible share. This ranking supports joint parameter tuning rather than one-factor optimization for safety-oriented design decisions.

For visual consistency, all quantitative analytics plots were regenerated using a unified serif typography and consistent axis-label conventions.

8 VALIDATION AND DISCUSSION

8.1 Empirical Consistency Check

PeopleFlow is evaluated as an engineering evaluation framework by checking consistency with established pedestrian-dynamics ranges rather than claiming exact one-to-one replication. Table 17 compares simulation trends with empirical benchmarks from classical and fire-engineering literature.

In evacuation literature, area-normalized pedestrian density of roughly 2–4 persons/m² is common in regular circulation, while 5–8 persons/m² indicates high-congestion conditions that often precede unstable flow Fruin (1971); Still (2014); Haghani (2020). Our area-normalized hotspot estimates from controlled validation runs fall in these bands during normal and stressed scenarios, respectively. The larger peak-density values in Table 8 are conservative kernel-level

hotspot intensity indicators used for ranking relative risk across layouts, not direct one-to-one field counts. Kernel-level density estimates represent localized interaction intensity and are used for relative comparison across layouts rather than direct real-world person-per-square-meter equivalence. These observations indicate that the simulation captures fundamental qualitative patterns observed in empirical pedestrian flow studies. Consistent with established evacuation trends, occupancy scaling produces nonlinear congestion growth (Table 16), wider exits reduce clearance time substantially (Table 14), and persistent bottlenecks emerge near narrow corridor transitions (Figures 15 and 16) in line with prior obstacle-form and group-seeking observations Yano (2018); Chen and Wang (2021).

8.2 Comparative Real-World Validation

In the absence of directly conducted evacuation drill data, we strengthen empirical grounding through two complementary approaches. First, PeopleFlow outputs are compared against published evacuation-time bands from controlled drill studies in matched setting categories (Table 18), showing full range overlap and mean absolute midpoint deviation of 0.83s. Second, trajectory-level behavioral consistency is verified against the public ETH/UCY pedestrian dataset (Table 19), confirming that simulated motion characteristics remain consistent with empirically observed crowd dynamics. Together these provide a two-layer validation strategy – scenario-level range alignment and trajectory-level behavioral consistency – appropriate for a comparative engineering framework where the primary contribution is relative safety ranking rather than exact trajectory prediction. Future work will incorporate controlled drill observations to further strengthen absolute calibration.

To strengthen practical credibility, we compare representative literature evacuation-time bands with PeopleFlow ranges under matched scenario families. Table 18 shows that PeopleFlow ranges are directionally aligned with published evacuation-study intervals while preserving the model's comparative-ranking scope Gwynne et al. (1999); Ronchi and Gwynne (2015); Ronchi and Nilsson (2013).

8.3 ETH/UCY Trajectory Consistency Benchmark

To complement range-based evacuation validation, we executed a trajectory-level consistency benchmark on public ETH/UCY scenes using the implemented PeopleFlow interaction model. Following established trajectory-benchmark practice Pellegrini et al. (2009); Alahi et al. (2016); Gupta et al. (2018); Salzmann et al. (2020); Mohamed et al. (2020), we evaluated two scenes (“seq_eth,” “seq_hotel”). Since ETH/UCY is a trajectory-forecasting benchmark rather than a full building-evacuation benchmark, we treat this experiment as external behavioral-consistency evidence, not as direct validation of full evacuation outcomes. We report six primary parameters: root-mean-square error (RMSE), average displacement error (ADE), final displacement error (FDE), oscillation strength, path smoothness, and successful-trajectory ratio. As summarized in Table 19, the aggregate RMSE is 0.628 with ADE 0.516 and FDE 0.542 over 13,953 trajectory samples; the same run reports oscillation strength of 15.60°, path smoothness of 0.913, successful-trajectory ratio of 0.977, and overlap proportion of 0.081. For methodological context, classical linear-velocity and basic social-force baselines in ETH/UCY literature generally show higher interaction-error sensitivity than interaction-aware models under dense crossing behavior; in safety-oriented evaluation settings, the successful-trajectory ratio of 0.977 is a primary robustness

indicator. The ETH/UCY benchmark demonstrates that simulated motion characteristics remain consistent with empirically observed pedestrian trajectory dynamics.

[‡]Trajectory-forecasting baselines in ETH/UCY literature optimize future-prediction objectives under protocol-specific observation/prediction windows. These are not directly comparable to PeopleFlow's full-scene replay-consistency setting, so the table provides methodological context instead of mixed-protocol numeric ranking.

These values indicate stable trajectory-level behavior under heterogeneous crowd scenes and provide an external quantitative sanity check complementary to scenario-level evacuation metrics. The overlap-proportion diagnostic is included as a conservative clipping-proxy measure for crowd proximity events and should be interpreted as a safety-oriented consistency signal rather than a standard ETH leaderboard metric.

8.4 Quantitative Agreement with Published Ranges

Beyond directional agreement, Table 20 reports midpoint-deviation metrics between published ranges and PeopleFlow outputs for matched setting categories.

The mean absolute midpoint deviation across these categories is 0.83 s, indicating close quantitative agreement for scenario-level engineering validation.

8.5 Interpretation for Safety Engineering

The key finding is layout sensitivity under stress. Identical behavioral parameters can produce large outcome shifts due to geometric topology and exit accessibility. For practitioners, this means that design alternatives should be compared through structured scenario matrices rather than a single nominal run. A focused Mall_Plan mini-sweep varying a group-cohesion parameter (weak-to-strong cohesion settings) showed consistent clearance delay and bottleneck persistence growth as cohesion increased, reinforcing the need to include social-group effects in safety-focused scenario studies.

The strongest practical use of PeopleFlow is comparative ranking: identifying which layout-policy pair reduces extreme congestion and improves clearance robustness. This aligns with modern risk-aware design practices and digital twin safety analytics Haghani (2020); Kuligowski (2013).

The proposed workflow enables architects and safety engineers to evaluate evacuation performance before construction or retrofitting. Rather than relying solely on compliance rules, designers can compare alternative layouts under structured stress scenarios and identify configurations that minimize congestion risk. The framework also supports integration into digital twin systems for continuous safety monitoring in smart buildings. Representative application domains include smart-building safety audits, fire evacuation planning, airport and rail-concourse safety design, emergency preparedness drills, and digital-twin what-if simulation. PeopleFlow can support safety certification workflows, smart building design, and evacuation planning in high-occupancy infrastructure.

This deployment pattern is particularly relevant for high-occupancy facilities such as transportation hubs, hospitals, and commercial centers, where transient surges in occupancy can quickly produce unsafe crowd states if geometric and routing constraints are not evaluated in advance. For multi-floor extension, the same routing logic can be lifted to a layered navigation graph in which each floor is a graph layer and stairs/elevators are modeled as inter-layer connectors

with vertical-transition penalties. This preserves the scenario-comparison workflow while extending applicability to 3D circulation networks.

8.6 Practical Deployment Workflow

For applied adoption, we recommend a four-step workflow that maps directly to safety-engineering practice:

1. Build a baseline model from the latest floor plan and verify extracted exits and obstacles.
2. Execute structured stress scenarios (high occupancy, blocked exits, routing-policy alternatives).
3. Compare design or policy variants using clearance, peak density, and bottleneck persistence metrics.
4. Integrate selected scenario sets into periodic digital-twin safety reviews for operational monitoring and retrofit planning.

This process converts simulation outputs into actionable planning evidence instead of one-off visual demonstrations. Session artifacts are persisted in a standardized JSON schema (run metadata, geometry hash, scenario parameters, per-step metrics, and aggregate outputs) to support direct ingestion by digital-twin orchestration services and building-management APIs. In online operation, live occupancy-counter streams can update local density estimates $\rho_e(t)$ in Eq. (5), enabling real-time route-cost adaptation under changing crowd conditions.

8.7 Reproducibility and Data Availability

The PeopleFlow codebase, scenario configurations, and reproducibility scripts are publicly maintained at <https://github.com/1206likith/PeopleFlow>. A persistent archival snapshot is available at <https://doi.org/10.5281/zenodo.19432157>. ETH/UCY benchmark inputs are sourced from the public ETH archive used in this study (https://data.vision.ee.ethz.ch/cvl/aem/ewap_dataset_full.tgz). Core paper artifacts can be regenerated through the backend command-line pipeline: from the repository root, enter “apps/backend” and run “python -m app.experiments.cli -paper-bundle,” then run “python -m app.experiments.cli -validate-eth -eth-download” to regenerate ETH outputs. For containerized execution, enter “infra” and run “docker compose -f docker-compose.prod.yml up -build.” All plots, tables, and ETH validation metrics can be regenerated using provided CLI pipeline. The multimedia supplement can be regenerated directly from the maintained script “python app/experiments/generate_multimedia_supplement.py,” ensuring that the narrated workflow video remains synchronized with the released analysis pipeline.

Generated outputs include aggregate paper tables/plots and ETH validation JSON/CSV artifacts under the research output folders, enabling independent reruns without manual post-processing. An interactive demo entry point is documented in the repository README under “Launch Demo” and can be started via the one-click launcher scripts.

8.8 Practical Design Implications

The results indicate that corridor-dominant geometries may require additional exit distribution to reduce bottleneck persistence. Open-hall layouts demonstrate more balanced exit utilization under congestion-aware routing strategies. These patterns provide directly actionable guidance for early-stage architectural and retrofit decisions.

8.9 Policy and Regulatory Implications

PeopleFlow is intended to complement, not replace, prescriptive compliance workflows. In practice, the framework can serve as an evidence layer for performance-based safety review by quantifying congestion hotspots, blocked-exit sensitivity, and clearance robustness under standardized stress scenarios. This supports transparent communication between designers, safety engineers, and authorities by adding reproducible scenario evidence to conventional rule-based checks, especially for high-occupancy facilities where geometry-induced bottlenecks can dominate risk. For jurisdictional review workflows, these outputs are suitable as supplementary technical evidence for Fire Marshal and Authority Having Jurisdiction (AHJ) discussions in performance-based evacuation assessments.

9 LIMITATIONS

This study has five main limitations. First, simulations are 2D and do not represent vertical circulation through stairs or elevators. A direct multi-floor extension would model each floor as a graph layer and encode stairs/elevators as inter-layer edges; the cost-function logic in Eq. (5) can then be evaluated on this augmented graph with additional vertical-transition penalties. Recent high-rise evacuation-elevator experiments further motivate this extension Mossberg et al. (2020). Second, hazard fields such as smoke and heat are not yet coupled to route choice in a physically calibrated way. Third, panic behavior is simplified, and human behavioral adaptation such as group cohesion and leader-following dynamics are approximated; group-level coordination and emergency communication effects are not explicitly modeled, which may influence evacuation dynamics in real deployments. Fourth, automatic floor-plan extraction can require correction for ambiguous drawings. Fifth, field-scale validation is currently limited to consistency checks against literature ranges rather than controlled human-subject evacuation trials. Supplementary airport/metro/office evaluations were expanded to full five-scenario coverage in this revision, but broader real-layout cohorts are still needed for stronger external generalization. These limitations define clear directions for future expansion and do not affect the comparative validity of the presented experiments.

10 FUTURE WORK

Future work will focus on (1) multi-floor 3D evacuation with staircase dynamics, (2) coupling with fire and smoke propagation models, (3) reinforcement-learning-assisted adaptive routing under hazards, and (4) live digital twin integration with sensor streams for online risk assessment. For advanced routing-policy research, future extensions can also incorporate preference-aware and compilation-based multi-agent pathfinding formulations to improve decision robustness in constrained networks Ho and Nakadai (2022); Surynek (2022). We also plan to expand the open benchmark set with additional plan categories such as hospital emergency floors, mall atria, railway concourses, and auditorium layouts for stronger cross-domain generalization, and to perform additional validation using controlled evacuation drills or publicly available trajectory datasets.

Practical Implications: For near-term deployment, the most actionable recommendation is to use PeopleFlow as a comparative screening tool during concept and retrofit stages: run structured stress scenarios early, rank alternatives by congestion-risk indicators and clearance robustness,

517 and forward only the strongest candidates into detailed engineering review. This workflow reduces
518 redesign latency and improves traceability of safety decisions.

11 CONCLUSION

519 This paper presented a reproducible geometry-aware multi-agent framework for structured
520 evaluation of building evacuation safety performance. Across a 450-run core matrix and
521 supplementary parametric/generalization experiments, we showed substantial sensitivity of
522 evacuation outcomes to geometric topology and routing assumptions, including an observed
523 8.27x range in clearance time and order-of-magnitude variation in peak density. Ablation,
524 parameter-sensitivity, and method-comparison experiments demonstrate that density-sensitive
525 path-selection policies and scenario design materially affect measured safety outcomes in structured
526 safety-evaluation settings.

527 The work advances evacuation research by framing simulation as an auditable, reproducible
528 framework rather than a one-off prediction tool. The resulting computational workflow supports
529 transparent comparison for architects, safety engineers, and smart-city planners who need evidence-
530 based prioritization of safer design and policy options. Overall, PeopleFlow provides a practical
531 foundation for experiment-driven evacuation safety evaluation, supports data-driven infrastructure
532 design, and serves as a pre-deployment evidence layer for standards-aligned review, design iteration,
533 and operational preparedness assessment.

534 PeopleFlow supports early-stage evacuation planning by enabling rapid evaluation of architectural
535 design alternatives prior to construction. The framework can assist safety engineers in identifying
536 congestion-prone regions, evaluating exit placement strategies, and testing emergency routing
537 policies within digital twin environments. Such simulation-driven design assessment can contribute
538 to safer high-occupancy infrastructure including transportation hubs, hospitals, and commercial
539 facilities. The workflow aligns with emerging digital twin methodologies for safety-aware
540 infrastructure modeling.

12 DIGITAL TWIN INTEGRATION SCHEMA AND API HOOKS

541 To support deployment-oriented digital-twin pipelines, PeopleFlow session artifacts follow a stable
542 schema that separates run identity, geometry state, scenario parameters, and outcome metrics.

543 Deployment-facing API hooks can expose this schema through standard endpoints:

- 544 • “/api/v2/simulations”
- 545 • “/api/v2/simulations/{id}/metrics”
- 546 • “/api/v2/simulations/{id}/artifacts”

547 In a live digital-twin loop, occupancy sensors update density estimates used in Eq. (5), while
548 archived artifacts provide audit-ready evidence for periodic safety review.

ACKNOWLEDGMENT

549 The authors used GitHub Copilot and ChatGPT for coding assistance and language polishing
550 during manuscript preparation. All AI-assisted outputs were reviewed, technically verified, and

revised by the authors before inclusion. The authors take full responsibility for the content, methodology, and technical accuracy of this work.

AUTHOR CONTRIBUTIONS

L.S. Kunchi conceived the PeopleFlow framework, designed and implemented the simulation pipeline, conducted all experiments, performed statistical analysis, and drafted the manuscript. [PhD Student Name] contributed to methodology validation, related work analysis, and critical manuscript revision. [Faculty Guide Name] supervised the research direction, provided technical guidance on pedestrian dynamics modeling and reproducibility design, and reviewed and approved the final manuscript. All authors have read and agreed to the published version of the manuscript.

REFERENCES

- Alahi, A., Goel, K., Ramanathan, V., Robicquet, A., Fei-Fei, L., and Savarese, S. (2016). Social lstm: Human trajectory prediction in crowded spaces. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 961–971. doi:10.1109/cvpr.2016.110
- Batty, M. (2012). Smart cities, big data. *Environment and Planning B: Planning and Design* 39, 191–193
- Chen, K., Song, X., and Ren, X. (2021). Modeling social interaction and intention for pedestrian trajectory prediction. *Physica A: Statistical Mechanics and its Applications* 570, 125790. doi:10.1016/j.physa.2021.125790
- Chen, S. and Wang, J. (2018). Agent-based modeling of pedestrian evacuation considering panic and guidance. *Simulation Modelling Practice and Theory* 82, 10–23
- Chen, X. and Wang, J. (2021). Entropy-based crowd evacuation modeling with seeking behavior of social groups. *IEEE Access* 9, 4653–4664. doi:10.1109/access.2020.3048054
- Ding, Y., Zhang, Y., and Huang, X. (2023). Intelligent emergency digital twin system for monitoring building fire evacuation. *Journal of Building Engineering* 77, 107416. doi:10.1016/j.job.2023.107416
- Duives, D. C., Daamen, W., and Hoogendoorn, S. P. (2013). State-of-the-art crowd motion simulation models. *Transportation Research Part C: Emerging Technologies* 37, 193–209
- Fruin, J. J. (1971). *Pedestrian Planning and Design* (New York, NY, USA: Metropolitan Association of Urban Designers and Environmental Planners)
- Gupta, A., Johnson, J., Fei-Fei, L., Savarese, S., and Alahi, A. (2018). Social gan: Socially acceptable trajectories with generative adversarial networks. In *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2255–2264. doi:10.1109/cvpr.2018.00240
- Gwynne, S., Galea, E., Owen, M., Lawrence, P., and Filippidis, L. (1999). A review of the methodologies used in the computer simulation of evacuation from the built environment. *Building and Environment* 34, 741–749
- Haghani, M. (2020). Empirical methods in pedestrian, crowd and panic dynamics: A review. *Safety Science* 129, 104841
- Helbing, D., Farkas, I., and Vicsek, T. (2000). Simulating dynamical features of escape panic. *Nature* 407, 487–490
- Helbing, D. and Molnar, P. (1995). Social force model for pedestrian dynamics. *Physical Review E* 51, 4282–4286

- Ho, F. and Nakadai, S. (2022). Preference-based multi-objective multi-agent path finding. *Autonomous Agents and Multi-Agent Systems* 37. doi:10.1007/s10458-022-09593-3
- Hurley, M. J., Gottuk, D., Hall, J. R. J., Harada, K., Kuligowski, E. D., Puchovsky, M., et al. (eds.) (2016). *SFPE Handbook of Fire Protection Engineering* (New York, NY, USA: Springer), 5 edn.
- Jahangir, M. F., Kamari, A., and Schultz, C. P. L. (2025). Unified and interoperable digital twin models for emergency evacuation using building simulation identity cards. *Automation in Construction* 175, 106225. doi:10.1016/j.autcon.2025.106225
- Jiang, L., Shi, J., Wang, C., and Pan, Z. (2023). Intelligent control of building fire protection system using digital twins and semantic web technologies. *Automation in Construction* 147, 104728. doi:10.1016/j.autcon.2022.104728
- Kirchner, A. and Schadschneider, A. (2002). Simulation of evacuation processes using a bionics-inspired cellular automaton model for pedestrian dynamics. *Physica A: Statistical Mechanics and its Applications* 312, 260–276
- Kuligowski, E. D. (2013). Predicting human behavior during fires. *Fire Technology* 49, 101–120
- Li, X., Liu, Y., Wang, K., and Wang, F.-Y. (2020). A recurrent attention and interaction model for pedestrian trajectory prediction. *IEEE/CAA Journal of Automatica Sinica* 7, 1361–1370. doi:10.1109/jas.2020.1003300
- Lopez-Carmona, M. A. and Paricio Garcia, A. (2022). Adaptive cell-based evacuation systems for leader-follower crowd evacuation. *Transportation Research Part C: Emerging Technologies* 140, 103699. doi:10.1016/j.trc.2022.103699
- Lovreglio, R., Gonzalez, V., Feng, Z., Amor, R., Spearpoint, M., Thomas, J., et al. (2020). Prototyping virtual reality serious games for building earthquake preparedness: The auckland city hospital case study. *Advanced Engineering Informatics* 32, 100900
- Mohamed, A., Qian, K., Elhoseiny, M., and Claudel, C. (2020). Social-stgcnn: A social spatio-temporal graph convolutional neural network for human trajectory prediction. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 14412–14420. doi:10.1109/cvpr42600.2020.01443
- Mossberg, A., Nilsson, D., and Andree, K. (2020). Unannounced evacuation experiment in a high-rise hotel building with evacuation elevators: A study of evacuation behaviour using eye-tracking. *Fire Technology* 57, 1259–1281. doi:10.1007/s10694-020-01046-1
- Nishinari, K., Kirchner, A., Namazi, A., and Schadschneider, A. (2004). Extended floor field CA model for evacuation dynamics. *IEICE Transactions on Information and Systems* E87-D, 726–732
- Osama, Z. (2024). The digital twin framework: A roadmap to the development of user-centred digital twin in the built environment. *Journal of Building Engineering* 98, 111081. doi:10.1016/j.jobbe.2024.111081
- Ozturk, G. B. (2021). Digital twin research in the AECO-FM industry. *Journal of Building Engineering* 40, 102730. doi:10.1016/j.jobbe.2021.102730
- Pan, X., Han, C. S., Dauber, K., and Law, K. H. (2007). A multi-agent based framework for the simulation of human and social behaviors during emergency evacuations. *AI & Society* 22, 113–132
- Pellegrini, S., Ess, A., Schindler, K., and van Gool, L. (2009). You'll never walk alone: Modeling social behavior for multi-target tracking. In *2009 IEEE 12th International Conference on Computer Vision*. 261–268. doi:10.1109/iccv.2009.5459260

- Ronchi, E. and Gwynne, S. (2015). Data collection for evacuation model validation. *Fire Safety Journal* 72, 137–151.
- Ronchi, E. and Nilsson, D. (2013). Fire evacuation in high-rise buildings: a review of human behaviour and modelling research. *Fire Science Reviews* 2, 1–21.
- Saleh, K. (2022). Pedestrian trajectory prediction for real-time autonomous systems via context-augmented transformer networks. *Sensors* 22, 7495. doi:10.3390/s22197495
- Salzmann, T., Ivanovic, B., Chakravarty, P., and Pavone, M. (2020). Trajectron++: Dynamically-feasible trajectory forecasting with heterogeneous data. In *Computer Vision – ECCV 2020* (Springer). 683–700. doi:10.1007/978-3-030-58523-5_40
- Still, G. K. (2014). *Introduction to Crowd Science* (Boca Raton, FL, USA: CRC Press)
- Sun, Q. and Turkan, Y. (2020). A BIM-based simulation framework for fire safety management and investigation of the critical factors affecting human evacuation performance. *Advanced Engineering Informatics* 44, 101093. doi:10.1016/j.aei.2020.101093
- Surynek, P. (2022). Problem compilation for multi-agent path finding: A survey. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence*. 5615–5622. doi:10.24963/ijcai.2022/783
- Wang, R., Song, X., Hu, Z., and Cui, Y. (2023). Spatio-temporal interaction aware and trajectory distribution aware graph convolution network for pedestrian multimodal trajectory prediction. *IEEE Transactions on Instrumentation and Measurement* 72, 1–11. doi:10.1109/tim.2022.3229733
- Wang, X., Chraibi, M., Chen, J., Li, R., and Ma, J. (2022). Modeling boundedly rational route choice in crowd evacuation processes. *Safety Science* 147, 105590. doi:10.1016/j.ssci.2021.105590
- Wehbe, R. and Shahrour, I. (2021). A BIM-based smart system for fire evacuation. *Future Internet* 13, 221. doi:10.3390/fi13090221
- Weidmann, U. (1993). *Transporttechnik der Fussgänger: Transporttechnische Eigenschaften des Fussgängerverkehrs* (Zurich, Switzerland: ETH Zurich, Institut für Verkehrsplanung, Transporttechnik, Strassen- und Eisenbahnbau)
- Yano, R. (2018). Effect of form of obstacle on speed of crowd evacuation. *Physical Review E* 97. doi:10.1103/physreve.97.032319
- Zhang, C., Lu, Y., and Tao, F. (2021). Digital twins for smart cities: A state-of-the-art review. *Engineering* 7, 1715–1730.
- Zheng, X., Zhong, T., and Liu, M. (2009). Modeling crowd evacuation of a building based on seven methodological approaches. *Building and Environment* 44, 437–445.
- Zhou, H., Ren, D., Xia, H., Fan, M., Yang, X., and Huang, H. (2021). Ast-gnn: An attention-based spatio-temporal graph neural network for interaction-aware pedestrian trajectory prediction. *Neurocomputing* 445, 298–308. doi:10.1016/j.neucom.2021.03.024

H: Metro Station Archetype




Figure 8. Example floor plan converted into simulation geometry with extracted walls and exits.




Figure 9. Simulation environment with agents navigating toward exits under dynamic routing.




Figure 10. Representative case-study floor plans used for multi-layout evacuation experiments.

Table 8. Comprehensive Results Across Nine Layouts and Five Scenarios (10 Runs per Configuration). Peak-density values are kernel-level hotspot intensity indicators for relative ranking and are not direct real-world persons/m² counts.

Case Study	Scenario	Evacuation Time (s)		Peak Density (agents/m ²)	
		Mean	Std	Mean	Std
Acad.	S1 Base	182.68	16.44	120.49	3.91
	S2	191.92	9.48	131.74	4.83
	HighOcc				
	S3	190.34	6.78	125.66	3.53
	Blocked				
	S4	38.26	2.21	24.55	2.30
	Routing				
	S5	140.14	32.15	104.46	12.85
Hosp.	Panic				
	S1 Base	30.12	2.08	8.54	0.99
	S2	32.64	2.45	12.92	1.64
	HighOcc				
	S3	46.96	3.49	18.37	2.35
	Blocked				
	S4	32.52	1.80	8.84	1.35
	Routing				
Hotel	S5	45.50	4.66	7.98	0.82
	Panic				
	S1 Base	138.02	31.96	100.13	15.16
	S2	121.02	27.47	80.16	15.75
	HighOcc				
	S3	153.24	39.78	110.32	11.44
	Blocked				
	S4	144.64	28.53	179.94	6.47
Libr.	Routing				
	S5	131.18	37.19	57.07	2.10
	Panic				
	S1 Base	52.84	2.97	21.80	2.00
	S2	57.22	4.35	33.99	3.17
	HighOcc				
	S3	59.08	3.79	30.25	1.55
	Blocked				
Mall	S4	58.76	5.23	22.26	1.09
	Routing				
	S5	70.96	10.73	20.93	2.54
	Panic				
	S1 Base	24.18	2.43	8.31	1.61
	S2	26.88	1.87	12.15	1.59
	HighOcc				
	S3	37.40	2.04	13.65	1.43
P1	Blocked				
	S4	35.10	1.67	9.30	1.45
	Routing				
	S5	42.86	4.63	8.66	1.21
	Panic				
	S1 Base	37.14	2.70	10.85	0.74
	S2	37.58	0.93	15.98	1.77
	HighOcc				
P2	S3	51.26	4.44	18.18	2.15
	Blocked				
	S4	39.02	5.32	10.92	2.51
	Routing				
	S5	46.58	3.88	9.40	1.21
	Panic				
	S1 Base	36.62	3.11	9.30	1.85
	S2	38.06	2.56	12.45	2.19
Frontiers	HighOcc				
	S3	39.52	2.63	12.91	2.92
	Blocked				

Table 9. Ablation Study Results

Configuration	Mean Time (s)	Mean Density (agents/m ²)
Full Model (S4)	7.24	3.96
No Congestion Routing	39.42	28.58
No Panic Variation	10.08	5.46

Table 10. Aggregate Method Summary (Benchmark Set)

Method	Clearance (s)	Peak Density (agents/m ²)	Completed (%)
Shortest path	301.96	140.52	86.05
Stochastic routing	189.23	70.44	72.45
Least-crowded routing	211.56	90.42	83.02
Guided variant	276.17	202.58	58.91

Table 11. Comparison with Publicly Used Baseline Assumptions

Method	Typical source	Evacuation-time behavior	Density/usage behavior
Shortest path	Widely used baseline Gwynne et al. (1999)	Often delayed by choke points	Higher local congestion near nearest exits
Random/stochastic routing	Agent-based baseline Pan et al. (2007)	More variable and less efficient exit usage	Reduced hotspot focus but weaker overall efficiency
Congestion-aware (PeopleFlow)	This work	More stable clearance under stress	Better exit balance with reduced persistent bottlenecks

Table 12. Blocked-Exit Stress Comparison (10 Runs, 120 Agents)

Policy	Mean time (s)	Peak density	Completion (%)
Nearest/shortest baseline	33.32 ± 2.39	21.20	95.0
Congestion-aware (PeopleFlow)	57.64 ± 3.74	13.79	95.0

Table 13. Supplementary Real-Layout Cohort (Expanded S1–S5, 10 Runs Each)

Layout	Scenario	Evacuation Time (s)	Peak Density (agents/m ²)
Airport terminal	S1.Baseline	34.26 ± 9.11	109.23 ± 34.02
	S2.HighOcc	29.20 ± 6.72	13.02 ± 3.68
	S3.Blocked	31.88 ± 6.44	108.14 ± 34.01
	S4.Routing	74.70 ± 56.04	46.06 ± 37.51
	S5.Panic	32.86 ± 3.63	31.97 ± 42.89
Metro station	S1.Baseline	31.52 ± 2.33	12.02 ± 1.71
	S2.HighOcc	32.10 ± 2.19	18.37 ± 2.02
	S3.Blocked	32.32 ± 2.57	12.78 ± 1.90
	S4.Routing	33.10 ± 3.82	9.67 ± 1.89
	S5.Panic	39.50 ± 5.56	10.92 ± 1.43
Office building	S1.Baseline	33.86 ± 5.07	7.86 ± 1.48
	S2.HighOcc	30.52 ± 3.35	8.83 ± 0.86
	S3.Blocked	34.38 ± 2.93	9.39 ± 1.81
	S4.Routing	27.80 ± 2.23	6.59 ± 0.93
	S5.Panic	34.08 ± 4.02	6.76 ± 1.29

Table 14. Controlled Comparison: Exit Width vs. Evacuation Time

Exit Width (m)	Mean Time (s)	Reduction vs. 1.0 m
1.0	141.2 ± 10.9	—
1.5	95.8 ± 8.6	32.2%
2.0	70.4 ± 7.3	50.1%

Table 15. Sensitivity Study: Desired Speed vs. Evacuation Outcomes

Desired Speed (m/s)	Mean Peak Density (agents/m ²)	
	Mean Time (s)	
1.0	126.4 ± 9.7	18.9
1.2	101.3 ± 8.5	22.1
1.5	81.2 ± 7.4	26.8

Table 16. Occupancy Scaling in Controlled Baseline Sweep

Occupancy (agents)	Mean Peak Density (agents/m ²)	
	Mean Time (s)	
50	62.5 ± 5.2	18.2 ± 1.6
100	101.3 ± 8.5	22.1 ± 2.1
150	147.8 ± 12.7	28.9 ± 2.8




Figure 11. Sensitivity heatmap of routing weights α and β (mean evacuation time in s).




Figure 12. Two-panel sensitivity summary: evacuation-time response to desired walking speed (left), and Sobol-style proxy total-order indices across occupancy, exit width, and desired speed (right).




Figure 13. Evacuation efficiency comparison across layouts and scenarios (mean evacuation time in s).




Figure 14. Routing-strategy comparison under controlled benchmarks (clearance in s, peak density in agents/m²).




Figure 15. Agent trajectories showing path adaptation under congestion-aware routing.

Table 17. Validation Against Empirical Benchmarks

Indicator	Literature Range	PeopleFlow Observation
Free-flow speed	1.2–1.4 m/s (Fruin, Weidmann)	Nominal target 1.2 m/s; low-density runs stay near target
Moderate density speed	Approx. 0.8–1.2 m/s	Speed reductions observed as local density rises
High density regime	Strong nonlinear slowdown / jams	High-density scenarios produce bottlenecks and delayed clearance
Density-flow trend	Unimodal flow behavior	Qualitative agreement in flow-density plots

Table 18. Comparative Real-World Validation (Evacuation Time)

Setting	Literature Range (s)	PeopleFlow Range (s)
Academic corridor drill	60–90	66–82
Retail floor segment	95–140	102–134
Transport concourse zone	150–230	168–214




Figure 16. Spatial heatmap of accumulated agent occupancy intensity over simulation frames; warmer colors indicate higher crowd concentration and recurrent congestion hotspots near constrained exits (kernel occupancy-intensity units).

Table 19. ETH/UCY Trajectory Consistency Benchmark Results

Scene	Samples	RMSE	ADE	FDE
seq_eth	8,188	0.527	0.432	0.465
seq_hotel	5,765	0.729	0.599	0.619
Aggregate	13,953	0.628	0.516	0.542
Benchmark context row [‡]	—	—	—	—

Table 20. Quantitative Agreement Metrics for Published Validation Ranges

Setting	Midpoint dev. (s)	Relative dev. (%)	Range overlap
Academic corridor drill	1.0	1.33	Full overlap
Retail floor segment	0.5	0.43	Full overlap
Transport concourse zone	1.0	0.53	Full overlap




Figure 17. Observed flow-density trend (density in agents/m², flow in agents/s) used for qualitative consistency checks.

Table 21. PeopleFlow Artifact Schema for Digital-Twin Integration

Schema field	Description
run_id	Unique execution identifier with timestamp and seed tag
geometry_hash	Deterministic hash of processed floor-plan geometry
scenario_config	Occupancy, routing policy, blocked-exit flags, and behavioral settings
step_metrics	Time-series of local density, flow, queue length, and evacuation progress
aggregate_metrics	Clearance time, peak density indicator, bottleneck duration, exit utilization
provenance	Software version, parameter file checksum, and execution environment




Figure 18. Run-to-run variability of evacuation performance across repeated simulations (clearance in s; peak density in agents/m²).




Figure 19. Runtime scalability profile across representative layouts with wall-clock mean $\pm$ std and completion-rate overlay (80 agents, 5 repeated runs).




Figure 20. Conceptual layered 3D routing extension with inter-floor connectors.